

# MCFSimplex Manual

version 1.00

date 29 - 08 - 2011

Alessandro Bertolini

Antonio Frangioni

Operations Research Group

Dipartimento di Informatica

Università di Pisa



# Chapter 1

## Manual

### 1.1 Introduction

This file is a short user manual for our distribution of MCFSimplex a C++ implementation of the Primal and Dual Network Simplex algorithm for Linear and Quadratic Min Cost Flow problems.

### 1.2 Standard Disclaimer

This code is provided "as is", without any explicit or implicit warranty that it will properly behave or it will suit you needs. Although codes reaching the distribution phase have usually been extensively tested, we cannot guarantee that they are absolutely bug-free (who can?). Any use of the codes is at you own risk: in no case we could be considered liable for any damage or loss you would eventually suffer, either directly or indirectly, for having used this code. More details about the non-warranty attached to this code are available in the license description file.

The code also comes with a "good will only" support: feel free to contact us for any comments/critics/bug report/request help you may have, we will be happy to try to answer and help you. But we cannot spend much time solving your problems, just the time to read a couple of e-mails and send you fast suggestions if the problem is easily solvable. Apart from that, we can't offer you any support.

### 1.3 License

This code is provided free of charge under "GNU Lesser General Public License": see the file doc/LGPL.txt.

### 1.4 Release

Current version of MCFSimplex is: 1.00

Current date is: August 29, 2011

## 1.5 How to Use It

This release comes out with the following files:

- Manual: This file.
- doc/LGPL.txt: Description of the LGPL license
- doc/refman.pdf: Pdf version of the full manual.
- doc/html/\*: Html version of the full manual.
- OPTUtils.h: Contains some small utilities such as random number generators and timing routines.
- MCFCClass.h: Contains the declaration of class MCFCClass; it is an *abstract* class with *pure virtual* methods, so that you cannot declare objects of type MCFCClass, but only of derived class of its. MCFClass offers a general interface for MCF solver codes, with methods for setting and reading the data of the problems, for solving it and retrieving solution informations, and so on. The actual Min Cost Flow solver distributed in this package conforms with the interface, i.e., it derives from MCFCClass; however, the idea is that applications using this interface would require almost no changes if any other of the available solvers implementing the interface is used. Carefully read the public interface of the class to understand how to use the public methods of the class. Also, some compile time switches let you choose important features of the algorithm, such as if the node “names” are the integers  $0 \dots n - 1$  rather than  $1 \dots n$ .
- MCFSimplex.h: Contains the declarations of class MCFSimplex. Being a derived class of MCFClass, most of its interface is defined and discussed in MCFCClass.h; however, some implementation dependent details (and compile-time switches) which are worth knowing are described in this file.
- MCFSimplex.C: Contains the implementation of the MCFSimplex class, which uses the Primal or Dual Network Simplex algorithm to solve the Min Cost Flow problem. You should not need to read it.

The following files are not a part of the actual MCF solver, they are only provided as an example of how to use it:

- Main.C: Contains an example of use of MCFSimplex; it reads the data from a file in DIMACS standard format, constructs the problem, solves it and prints the results.
- Config.txt: Allows (if exists) to set some important parameters for the solver.
- sample.dmx: An example of a tiny (4 nodes, 4 arcs) network in DIMACS standard format; the comments in the file also describe the format itself.
- sample.qdmx: An example of tiny (4 nodes, 4 arcs) quadratic network in DIMACS extended format; the comments in the file also describe the format itself.
- makefile: A makefile of the example.

## 1.6 Testing the distribution

To compile the example, edit the makefile and do the proper changes (such as the compiler name), then just type “make”. To execute the example:

```
MCFSolve sample.dmx
```

The output should be something like

Optimal Objective Function value = 9

Solution time (s): user 0, system 0



# Contents

<b>1</b>	<b>Manual</b>	<b>i</b>
1.1	Introduction . . . . .	i
1.2	Standard Disclaimer . . . . .	i
1.3	License . . . . .	i
1.4	Release . . . . .	i
1.5	How to Use It . . . . .	ii
1.6	Testing the distribution . . . . .	ii
<b>2</b>	<b>Module Index</b>	<b>1</b>
2.1	Modules . . . . .	1
<b>3</b>	<b>Class Index</b>	<b>3</b>
3.1	Class Hierarchy . . . . .	3
<b>4</b>	<b>Class Index</b>	<b>5</b>
4.1	Class List . . . . .	5
<b>5</b>	<b>File Index</b>	<b>7</b>
5.1	File List . . . . .	7
<b>6</b>	<b>Module Documentation</b>	<b>9</b>
6.1	Compile-time switches in MCFCClass.h . . . . .	9
6.1.1	Detailed Description . . . . .	9
6.1.2	Define Documentation . . . . .	9
6.1.2.1	USERNAME0 . . . . .	9
6.2	Classes in MCFCClass.h . . . . .	10
6.3	Compile-time switches in MCFSimplex.h . . . . .	11
6.3.1	Detailed Description . . . . .	11
6.3.2	Define Documentation . . . . .	11
6.3.2.1	QUADRATICCOST . . . . .	11

6.4	Classes in MCFSimplex.h . . . . .	12
6.5	Compile-time switches in OPTutils.h . . . . .	13
6.5.1	Detailed Description . . . . .	13
6.5.2	Define Documentation . . . . .	13
6.5.2.1	OPT_RANDOM . . . . .	13
6.5.2.2	OPT_TIMERS . . . . .	13
6.5.2.3	OPT_USE_NAMESPACES . . . . .	14
6.6	Classes in OPTutils.h . . . . .	15
6.7	Functions in OPTutils.h . . . . .	16
6.7.1	Function Documentation . . . . .	16
6.7.1.1	DfltDfInpt . . . . .	16
<b>7</b>	<b>Class Documentation</b>	<b>17</b>
7.1	MCFCClass::Eps< T > Class Template Reference . . . . .	17
7.1.1	Detailed Description . . . . .	17
7.2	MCFCClass::Inf< T > Class Template Reference . . . . .	18
7.2.1	Detailed Description . . . . .	18
7.3	MCFCClass Class Reference . . . . .	19
7.3.1	Detailed Description . . . . .	25
7.3.2	Member Typedef Documentation . . . . .	26
7.3.2.1	FONumber . . . . .	26
7.3.3	Member Enumeration Documentation . . . . .	26
7.3.3.1	MCFAnswer . . . . .	26
7.3.3.2	MCFFlFrmmt . . . . .	26
7.3.3.3	MCFParam . . . . .	26
7.3.3.4	MCFStatus . . . . .	27
7.3.4	Constructor & Destructor Documentation . . . . .	27
7.3.4.1	MCFCClass . . . . .	27
7.3.4.2	~MCFCClass . . . . .	27
7.3.5	Member Function Documentation . . . . .	28
7.3.5.1	AddArc . . . . .	28
7.3.5.2	AddNode . . . . .	28
7.3.5.3	ChangeArc . . . . .	28
7.3.5.4	CheckDSol . . . . .	28
7.3.5.5	CheckPSol . . . . .	28
7.3.5.6	ChgCost . . . . .	28
7.3.5.7	ChgCosts . . . . .	29



7.3.5.8	ChgDfct	29
7.3.5.9	ChgDfcts	29
7.3.5.10	ChgQCoef	30
7.3.5.11	ChgQCoef	30
7.3.5.12	ChgUCap	30
7.3.5.13	ChgUCaps	31
7.3.5.14	CloseArc	31
7.3.5.15	DelArc	31
7.3.5.16	DelNode	31
7.3.5.17	ETZ	31
7.3.5.18	GetPar	32
7.3.5.19	GetPar	32
7.3.5.20	GEZ	32
7.3.5.21	GT	32
7.3.5.22	GTZ	32
7.3.5.23	HaveNewPi	32
7.3.5.24	HaveNewX	33
7.3.5.25	IsClosedArc	33
7.3.5.26	IsDeletedArc	33
7.3.5.27	LEZ	33
7.3.5.28	LoadDMX	33
7.3.5.29	LoadNet	34
7.3.5.30	LT	35
7.3.5.31	LTZ	35
7.3.5.32	MCFArcs	35
7.3.5.33	MCFCost	35
7.3.5.34	MCFCosts	35
7.3.5.35	MCFCosts	36
7.3.5.36	MCFDfct	36
7.3.5.37	MCFDfcts	36
7.3.5.38	MCFDfcts	36
7.3.5.39	MCFENde	37
7.3.5.40	MCFENdes	37
7.3.5.41	MCFGetDFO	37
7.3.5.42	MCFGetFO	37
7.3.5.43	MCFGetPi	37

7.3.5.44	MCFGetPi	38
7.3.5.45	MCFGetRC	38
7.3.5.46	MCFGetRC	38
7.3.5.47	MCFGetRC	38
7.3.5.48	MCFGetState	39
7.3.5.49	MCFGetStatus	39
7.3.5.50	MCFGetUnbCycl	39
7.3.5.51	MCFGetUnfCut	40
7.3.5.52	MCFGetX	40
7.3.5.53	MCFGetX	40
7.3.5.54	MCFm	40
7.3.5.55	MCFmmax	41
7.3.5.56	MCFn	41
7.3.5.57	MCFnmax	41
7.3.5.58	MCFPutState	41
7.3.5.59	MCFQCoef	41
7.3.5.60	MCFQCoef	42
7.3.5.61	MCFQCoef	42
7.3.5.62	MCFSNde	42
7.3.5.63	MCFSNdes	42
7.3.5.64	MCFUCap	42
7.3.5.65	MCFUCaps	42
7.3.5.66	MCFUCaps	43
7.3.5.67	OpenArc	43
7.3.5.68	PreProcess	43
7.3.5.69	SetMCFTime	43
7.3.5.70	SetPar	43
7.3.5.71	SetPar	44
7.3.5.72	SolveMCF	45
7.3.5.73	TimeMCF	45
7.3.5.74	TimeMCF	45
7.3.5.75	WriteMCF	45
7.3.6	Member Data Documentation	45
7.3.6.1	MaxIter	45
7.3.6.2	MaxTime	46
7.3.6.3	status	46

---

7.4	MCFCClass::MCFException Class Reference	47
7.4.1	Detailed Description	47
7.5	MCFSimplex Class Reference	48
7.5.1	Detailed Description	49
7.5.2	Member Enumeration Documentation	49
7.5.2.1	enumPrngRl	49
7.5.2.2	SimplexParam	50
7.5.3	Constructor & Destructor Documentation	50
7.5.3.1	MCFSimplex	50
7.5.4	Member Function Documentation	50
7.5.4.1	AddArc	50
7.5.4.2	AddNode	50
7.5.4.3	ChangeArc	50
7.5.4.4	CloseArc	51
7.5.4.5	DelArc	51
7.5.4.6	DelNode	51
7.5.4.7	IsClosedArc	51
7.5.4.8	IsDeletedArc	51
7.5.4.9	LoadNet	51
7.5.4.10	MCFCost	52
7.5.4.11	MCFDfct	52
7.5.4.12	MCFENde	52
7.5.4.13	MCFGetFO	52
7.5.4.14	MCFGetRC	52
7.5.4.15	MCFSNde	53
7.5.4.16	MCFUCap	53
7.5.4.17	OpenArc	53
7.5.4.18	SetAlg	53
7.5.4.19	SetPar	53
7.5.4.20	SetPar	54
7.5.4.21	SolveMCF	54
7.6	MCFCClass::MCFState Class Reference	55
7.6.1	Detailed Description	55
7.7	OPTrand Class Reference	56
7.7.1	Detailed Description	56
7.7.2	Member Function Documentation	56

7.7.2.1	rand	56
<b>8</b>	<b>File Documentation</b>	<b>57</b>
8.1	Main.C File Reference	57
8.1.1	Detailed Description	57
8.2	MCFCClass.h File Reference	59
8.2.1	Detailed Description	59
8.3	MCFSimplex.h File Reference	60
8.3.1	Detailed Description	60
8.4	OPTUtils.h File Reference	61
8.4.1	Detailed Description	61

## Chapter 2

# Module Index

### 2.1 Modules

Here is a list of all modules:

Compile-time switches in MCFCClass.h . . . . .	9
Classes in MCFCClass.h . . . . .	10
Compile-time switches in MCFSimplex.h . . . . .	11
Classes in MCFSimplex.h . . . . .	12
Compile-time switches in OPTutils.h . . . . .	13
Classes in OPTutils.h . . . . .	15
Functions in OPTutils.h . . . . .	16



# Chapter 3

## Class Index

### 3.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

MCFCClass::Eps< T > . . . . .	17
MCFCClass::Inf< T > . . . . .	18
MCFCClass . . . . .	19
MCFSimplex . . . . .	48
MCFCClass::MCFCException . . . . .	47
MCFCClass::MCFCState . . . . .	55
OPTrand . . . . .	56





# Chapter 4

## Class Index

### 4.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

<a href="#">MCFCClass::Eps&lt; T &gt;</a>	17
<a href="#">MCFCClass::Inf&lt; T &gt;</a>	18
<a href="#">MCFCClass</a>	19
<a href="#">MCFCClass::MCFException</a>	47
<a href="#">MCFSimplex</a>	48
<a href="#">MCFCClass::MCFState</a>	55
<a href="#">OPTrand</a>	56



# Chapter 5

## File Index

### 5.1 File List

Here is a list of all documented files with brief descriptions:

<a href="#">Main.C</a>	57
<a href="#">MCFCClass.h</a>	59
<a href="#">MCFSimplex.h</a>	60
<a href="#">OPTUtils.h</a>	61



## Chapter 6

# Module Documentation

### 6.1 Compile-time switches in MCFClass.h

#### Defines

- `#define USERNAME0 0`

#### 6.1.1 Detailed Description

These macros control some important details of the class interface. Although using macros for activating features of the interface is not very C++, switching off some unused features may allow some implementation to be more efficient in running time or memory.

#### 6.1.2 Define Documentation

##### 6.1.2.1 `#define USERNAME0 0`

Decides if 0 or 1 is the "name" of the first node. If `USERNAME0 == 1`, (warning: it has to be *\*exactly\** 1), then the node names go from 0 to `n - 1`, otherwise from 1 to `n`. Note that this does not affect the position of the deficit in the deficit vectors, i.e., the deficit of the `i`-th node - be its "name" `'i'` or `'i - 1'` - is always in the `i`-th position of the vector.

## 6.2 Classes in MCFClass.h

### Classes

- class [MCFClass](#)

## 6.3 Compile-time switches in MCFSimplex.h

### Defines

- `#define QUADRATICCOST 0`

#### 6.3.1 Detailed Description

There is only one macro in [MCFSimplex](#), but it is very important !

#### 6.3.2 Define Documentation

##### 6.3.2.1 `#define QUADRATICCOST 0`

Setting `QUADRATICCOST = 1` the solver can solve problems with linear and quadratic costs too, but it can use only the Primal Simplex algorithm. The solver uses "arcType" struct with the additional field "quadraticCost" for the quadratic coefficients. The field "ident" isn't loaded because the solver doesn't use the tripartition TLU. To determine close arcs and deleted arcs the solver uses this convention:

- close arcs have the field "cost" to INFINITY (`Inf<FNumber>()` in [MCFClass.h](#))
- deleted arcs have the field "upper" to INFINITY, and the field "tail" and "head" point to NULL. Then the solver loads the variables "ignoredEnteringArc" and "firstIgnoredEnteringArc", used to avoid nasty loops during the execution of the Quadratic Primal Simplex algorithm. Instead, if `QUADRATICCOST = 0` the solver can solve only problems with linear costs. The field "quadraticCost" would be useless and it isn't be loaded. Primal Simplex and Dual Simplex use the tripartition TLU to divide the arcs. So the solver loads the field "ident", which differentiates the set of the arcs in: deleted arcs, closed arcs, arcs in T, arcs in L, arcs in U. With `QUADRATICCOST = 0` the solver cannot solve problems with quadratic costs, but it solve the problems with linear costs more quickly.

## 6.4 Classes in MCFSimplex.h

### Classes

- class [MCFSimplex](#)



## 6.5 Compile-time switches in OPTutils.h

### Defines

- #define `OPT_USE_NAMESPACES` 0
- #define `OPT_TIMERS` 5
- #define `OPT_RANDOM` 1

### 6.5.1 Detailed Description

These macros control how the classes OPTTimers and `OPTrand` are implemented; choose the appropriate value for your environment, or program a new version if no value suits you. Also, namespaces can be eliminated if they create problems.

### 6.5.2 Define Documentation

#### 6.5.2.1 #define OPT\_RANDOM 1

The class `OPTrand` is defined below to give an abstract interface to the different random generators that are used in different platforms. This is needed since random generators are one of the less standard parts of the C++ library. The value of the `OPT_RANDOM` constant selects among the different timing routines:

- 0 = an hand-made implementation of a rather good random number generator is used; note that this assumes that long ints  $\geq 32$  bits
- 1 = standard `rand()` / `srand()` pair, common to all C libraries but not very sophisticated
- 2 = `drand48()` / `srand48()`, common on Unix architectures and pretty good.

Any unsupported value would simply make the functions to report constant zero, which is not nice but useful to quickly fix problems if you don't use random numbers at all.

#### 6.5.2.2 #define OPT\_TIMERS 5

The class OPTtimers is defined below to give an abstract interface to the different timing routines that are used in different platforms. This is needed since time-related functions are one of the less standard parts of the C++ library. The value of the `OPT_TIMERS` constant selects among the different timing routines:

- 1 = Use the Unix `times()` routine in `sys/times.h`
- 2 = As 1 but uses `sys/timeb.h` (typical for Microsoft(TM) compilers)
- 3 = Still use `times()` of `sys/times.h`, but return wallclock time rather than CPU time
- 4 = As 3 but uses `sys/timeb.h` (typical for Microsoft(TM) compilers)
- 5 = return the user time obtained with ANSI C `clock()` function; this may result in more accurate running times w.r.t. but may be limited to  $\sim 72$  hours on systems where ints are 32bits.

- 6 = Use the Unix `gettimeofday()` routine of `sys/time.h`.

Any unsupported value would simply make the class to report constant zero as the time.

The values 1 .. 4 rely on the constant `CLK_TCK` for converting between clock ticks and seconds; for the case where the constant is not defined by the compiler -- should not happen, but it does -- or it is defined in a wrong way, the constant is re-defined below.

#### **6.5.2.3 `#define OPT_USE_NAMESPACES 0`**

Setting `OPT_USE_NAMESPACES == 0` should instruct all codes that use OPTutils stuff to avoid using namespaces; to start with, the common namespace `OPTutils_di_unipi_it`, that contains all the types defined herein, is *\*not\** defined.

## 6.6 Classes in OPTutils.h

### Classes

- class [OPTrand](#)

## 6.7 Functions in OPTutils.h

### Functions

- `template<class T >`  
`void DfltDSfInpt (istream *iStrm, T &Param, const T Dflt, const char cmntc= '#')`

#### 6.7.1 Function Documentation

**6.7.1.1** `template<class T > void DfltDSfInpt (istream * iStrm, T & Param, const T Dflt, const char cmntc = '#') [inline]`

Template function for reading parameters from a istream. The function is "safe" because it works also if the istream is not given, is not be long enough or contains erroneous things.

Given a &istream (possibly NULL), `DfltDSfInpt()` attempts to read Param out of it, skipping any line that begins with the comment carachter (defaulted to '#'), any blank line and any line starting with anything that can not be interpreted as a 'T'. If, for any reason, the read operation fails, then the parameter is given the default value 'Dflt'. Otherwise, all the rest of the line up to the nearest newline ('

') carachter is flushed.

## Chapter 7

# Class Documentation

### 7.1 MCFClass::Eps< T > Class Template Reference

```
#include <MCFClass.h>
```

#### Public Member Functions

- `operator T ()`

#### 7.1.1 Detailed Description

**template<typename T> class MCFClass::Eps< T >**

Very small class to simplify extracting the "machine epsilon" for a basic type (FNumber, CNumber); just use `Eps<type>()`.

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

## 7.2 MCFClass::Inf< T > Class Template Reference

```
#include <MCFClass.h>
```

### Public Member Functions

- `operator T ()`

### 7.2.1 Detailed Description

**template<typename T> class MCFClass::Inf< T >**

Very small class to simplify extracting the "+ infinity" value for a basic type (FNumber, CNumber, Index); just use `Inf<type>()`.


The documentation for this class was generated from the following file:

- [MCFClass.h](#)

## 7.3 MCFCClass Class Reference

```
#include <MCFCClass.h>
```

Inheritance diagram for MCFCClass:



### Classes

- class [Eps](#)
- class [Inf](#)
- class [MCFException](#)
- class [MCFCState](#)

### Public Types

#### Public types

The [MCFCClass](#) defines four main public types:

- *Index*, the type of arc and node indices;
- *FNumber*, the type of flow variables, arc capacities, and node deficits;
- *CNumber*, the type of flow costs, node potentials, and arc reduced costs;
- *FONumber*, the type of objective function value.

By re-defining the types in this section, most *MCFSolver* should be made to work with any reasonable choice of data type (= one that is capable of properly representing the data of the instances to be solved). This may be relevant due to an important property of MCF problems: *\*if all arc capacities and node deficits are integer, then there exists an integral optimal primal solution\**, and *\*if all arc costs are integer, then there exists an integral optimal dual solution\**. Even more importantly, *\*many solution algorithms will in fact produce an integral primal/dual solution for free\**, because *\*every primal/dual solution they generate during the solution process is naturally integral\**. Therefore, one can use integer data types to represent everything connected with flows and/or costs if the corresponding data is integer in all instances one needs to solve. This directly translates in significant memory savings and/or speed improvements.

It is the user's responsibility to ensure that these types are set to reasonable values\*. So, the experienced user may want to experiment with setting this data properly if memory footprint and/or speed is a primary concern. Note, however, that *\*not all solution algorithms will happily accept integer data\**; one example are Interior-Point approaches, which require both flow and cost variables to be continuous (float). So, the viability of setting integer data (as well as its impact on performances) is strictly related to the specific kind of algorithm used. Since these types are common to all derived classes, they have to be set taking into account the needs of all the solvers that are going to be used, and adapting to the "worst case"; of course, `FNumber == CNumber == double` is going to always be an acceptable "worst case" setting. [MCFCClass](#) may in a future be defined as a template class, with these as template parameters, but this is currently deemed overkill and avoided.

Finally, note that the above integrality property only holds for *\*linear\** MCF problems. If any arc has a nonzero quadratic cost coefficient, optimal flows and potentials may be fractional even if all the data of the problem (comprised quadratic cost coefficients) is integer. Hence, for *\*quadratic\** MCF solvers, a setting like `FNumber == CNumber == double` is actually mandatory\*, for any reasonable algorithm will typically misbehave otherwise.

- enum `MCFParam` {  
`kMaxTime` = 0, `kMaxIter`, `kEpsFlw`, `kEpsDfct`,  
`kEpsCst`, `kReopt`, `kLastParam` }
- enum `MCFStatus` {  
`kUnSolved` = -1, `kOK` = 0, `kStopped`, `kUnfeasible`,  
`kUnbounded`, `kError` }
- enum `MCFAnswer` { `kNo` = 0, `kYes` }
- enum `MCFFIFrmt` { `kDimacs` = 0, `kQDimacs`, `kMPS`, `kFWMPS` }
- typedef unsigned int `Index`  
*index of a node or arc ( >= 0 )*
- typedef `Index * Index_Set`  
*set (array) of indices*
- typedef const `Index cIndex`  
*a read-only index*
- typedef `cIndex * cIndex_Set`  
*read-only index array*
- typedef int `SIndex`  
*index of a node or arc*
- typedef `SIndex * SIndex_Set`  
*set (array) of indices*
- typedef const `SIndex cSIndex`  
*a read-only index*
- typedef `cSIndex * cSIndex_Set`  
*read-only index array*
- typedef double `FNumber`  
*type of arc flow*
- typedef `FNumber * FRow`  
*vector of flows*
- typedef const `FNumber cFNumber`  
*a read-only flow*
- typedef `cFNumber * cFRow`  
*read-only flow array*
- typedef double `CNumber`  
*type of arc flow cost*



- typedef [CNumber](#) \* [CRow](#)  
*vector of costs*
- typedef const [CNumber](#) [cCNumber](#)  
*a read-only cost*
- typedef [cCNumber](#) \* [cCRow](#)  
*read-only cost array*
- typedef double [FONumber](#)
- typedef const [FONumber](#) [cFONumber](#)  
*a read-only o.f. value*
- typedef [MCFState](#) \* [MCFStatePtr](#)  
*pointer to a [MCFState](#)*

## Public Member Functions

### Constructors

- [MCFClass](#) ([cIndex](#) nmx=0, [cIndex](#) mmx=0)

### Other initializations

- virtual void [LoadNet](#) ([cIndex](#) nmx=0, [cIndex](#) mmx=0, [cIndex](#) pn=0, [cIndex](#) pm=0, [cFRow](#) pU=NULL, [cCRow](#) pC=NULL, [cFRow](#) pDfct=NULL, [cIndex\\_Set](#) pSn=NULL, [cIndex\\_Set](#) pEn=NULL)=0
- virtual void [LoadDMX](#) (istream &DMXs, bool IsQuad=false)
- virtual void [PreProcess](#) (void)
- virtual void [SetPar](#) (int par, int val)
- virtual void [SetPar](#) (int par, double val)
- virtual void [GetPar](#) (int par, int &val)
- virtual void [GetPar](#) (int par, double &val)
- virtual void [SetMCFTIME](#) (bool TimeIt=true)

### Solving the problem

- virtual void [SolveMCF](#) (void)=0
- int [MCFGetStatus](#) (void)

### Reading flow solution

- virtual void [MCFGetX](#) ([FRow](#) F, [Index\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [cFRow](#) [MCFGetX](#) (void)
- virtual bool [HaveNewX](#) (void)

### Reading potentials

- virtual void [MCFGetPi](#) ([CRow](#) P, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [cCRow](#) [MCFGetPi](#) (void)
- virtual bool [HaveNewPi](#) (void)

**Reading reduced costs**

- virtual void [MCFGetRC](#) ([CRow](#) CR, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [cCRow](#) [MCFGetRC](#) (void)
- virtual [CNumber](#) [MCFGetRC](#) ([cIndex](#) i)=0

**Reading the objective function value**

- virtual [FONumber](#) [MCFGetFO](#) (void)=0
- virtual [FONumber](#) [MCFGetDFO](#) (void)

**Getting unfeasibility certificate**

- virtual [FNumber](#) [MCFGetUnfCut](#) ([Index\\_Set](#) Cut)

**Getting unboundedness certificate**

- virtual [Index](#) [MCFGetUnbCycl](#) ([Index\\_Set](#) Pred, [Index\\_Set](#) ArcPred)

**Saving/restoring the state of the solver**

- virtual [MCFStatePtr](#) [MCFGetState](#) (void)
- virtual void [MCFPutState](#) ([MCFStatePtr](#) S)

**Time the code**

- void [TimeMCF](#) (double &t\_us, double &t\_ss)
- double [TimeMCF](#) (void)

**Check the solutions**

- void [CheckPSol](#) (void)
- void [CheckDSol](#) (void)

**Reading graph size**

- [Index](#) [MCFnmax](#) (void)
- [Index](#) [MCFmmax](#) (void)
- [Index](#) [MCFn](#) (void)
- [Index](#) [MCFm](#) (void)

**Reading graph topology**

- virtual void [MCFArcs](#) ([Index\\_Set](#) Startv, [Index\\_Set](#) Endv, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [Index](#) [MCFSNde](#) ([cIndex](#) i)=0
- virtual [Index](#) [MCFENde](#) ([cIndex](#) i)=0
- virtual [cIndex\\_Set](#) [MCFSNdes](#) (void)
- virtual [cIndex\\_Set](#) [MCFENdes](#) (void)

**Reading arc costs**

- virtual void [MCFCosts](#) ([CRow](#) Costv, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0

- virtual [CNumber MCFCost](#) ([cIndex](#) i)=0
- virtual [cCRow MCFCosts](#) (void)
- virtual void [MCFQCoef](#) ([CRow](#) Qv, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- virtual [CNumber MCFQCoef](#) ([cIndex](#) i)
- virtual [cCRow MCFQCoef](#) (void)

#### Reading arc capacities

- virtual void [MCFUCaps](#) ([FRow](#) UCapv, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [FNumber MCFUCap](#) ([cIndex](#) i)=0
- virtual [cFRow MCFUCaps](#) (void)

#### Reading node deficits

- virtual void [MCFDfcts](#) ([FRow](#) Dfctv, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [FNumber MCFDfct](#) ([cIndex](#) i)=0
- virtual [cFRow MCFDfcts](#) (void)

#### Write problem to file

- virtual void [WriteMCF](#) (ostream &oStrm, int frmt=0)

#### Changing the costs

- virtual void [ChgCosts](#) ([cCRow](#) NCost, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgCost](#) ([Index](#) arc, [CNumber](#) NCost)=0
- virtual void [ChgQCoef](#) ([cCRow](#) NQCoef=NULL, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- virtual void [ChgQCoef](#) ([Index](#) arc, [CNumber](#) NQCoef)

#### Changing the capacities

- virtual void [ChgUCaps](#) ([cFRow](#) NCap, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgUCap](#) ([Index](#) arc, [cFNumber](#) NCap)=0

#### Changing the deficits

- virtual void [ChgDfcts](#) ([cFRow](#) NDfct, [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgDfct](#) ([Index](#) node, [cFNumber](#) NDfct)=0

#### Changing graph topology

- virtual void [CloseArc](#) ([cIndex](#) name)=0
- virtual bool [IsClosedArc](#) ([cIndex](#) name)=0
- virtual void [DelNode](#) ([cIndex](#) name)=0
- virtual void [OpenArc](#) ([cIndex](#) name)=0
- virtual [Index AddNode](#) ([cFNumber](#) aDfct)=0
- virtual void [ChangeArc](#) ([cIndex](#) name, [cIndex](#) nSN=Inf< [Index](#) >(), [cIndex](#) nEN=Inf< [Index](#) >())=0

- virtual void [DelArc](#) ([cIndex](#) name)=0
- virtual bool [IsDeletedArc](#) ([cIndex](#) name)=0
- virtual [Index AddArc](#) ([cIndex](#) Start, [cIndex](#) End, [cFNumber](#) aU, [cCNumber](#) aC)=0

### Destructor

- virtual [~MCFClass](#) ()

## Protected Member Functions

### Managing comparisons.

*The following methods are provided for making it easier to perform comparisons, with and without tolerances.*

- template<class T >  
bool [ETZ](#) (T x, const T eps)
- template<class T >  
bool [GTZ](#) (T x, const T eps)
- template<class T >  
bool [GEZ](#) (T x, const T eps)
- template<class T >  
bool [LTZ](#) (T x, const T eps)
- template<class T >  
bool [LEZ](#) (T x, const T eps)
- template<class T >  
bool [GT](#) (T x, T y, const T eps)
- template<class T >  
bool [LT](#) (T x, T y, const T eps)

## Protected Attributes

- [Index n](#)  
*total number of nodes*
- [Index nmax](#)  
*maximum number of nodes*
- [Index m](#)  
*total number of arcs*
- [Index mmax](#)  
*maximum number of arcs*
- int [status](#)
- bool [Senstv](#)  
*true <=> the latest optimal solution should be exploited*
- OPTtimers \* [MCFt](#)  
*timer for performances evaluation*
- [FNumber EpsFlw](#)

*precision for comparing arc flows / capacities*

- [FNumber EpsDfct](#)

*precision for comparing node deficits*

- [CNumber EpsCst](#)

*precision for comparing arc costs*

- double [MaxTime](#)
- int [MaxIter](#)

### 7.3.1 Detailed Description

This abstract base class defines a standard interface for (linear or convex quadratic separable) Min Cost Flow (MCF) problem solvers.

The data of the problem consist of a (directed) graph  $G = (N, A)$  with  $n = |N|$  nodes and  $m = |A|$  (directed) arcs. Each node 'i' has a deficit  $b[i]$ , i.e., the amount of flow that is produced/consumed by the node: source nodes (which produce flow) have negative deficits and sink nodes (which consume flow) have positive deficits. Each arc '(i, j)' has an upper capacity  $U[i, j]$ , a linear cost coefficient  $C[i, j]$  and a (non negative) quadratic cost coefficient  $Q[i, j]$ . Flow variables  $X[i, j]$  represents the amount of flow to be sent on arc (i, j). Parallel arcs, i.e., multiple copies of the same arc '(i, j)' (with possibly different costs and/or capacities) are in general allowed. The formulation of the problem is therefore:

$$\begin{aligned} \min \quad & \sum_{(i,j) \in A} C[i, j]X[i, j] + Q[i, j]X[i, j]^2/2 \\ (1) \quad & \sum_{(j,i) \in A} X[j, i] - \sum_{(i,j) \in A} X[i, j] = b[i] \quad i \in N \\ (2) \quad & 0 \leq X[i, j] \leq U[i, j] \quad (i, j) \in A \end{aligned}$$

The  $n$  equations (1) are the flow conservation constraints and the  $2m$  inequalities (2) are the flow non-negativity and capacity constraints. At least one of the flow conservation constraints is redundant, as the demands must be balanced ( $\sum_{i \in N} b[i] = 0$ ); indeed, exactly  $n - \text{ConnectedComponents}(G)$  flow conservation constraints are redundant, as demands must be balanced in each connected component of  $G$ . Let us denote by  $QA$  and  $LA$  the disjoint subsets of  $A$  containing, respectively, "quadratic" arcs (with  $Q[i, j] > 0$ ) and "linear" arcs (with  $Q[i, j] = 0$ ); the (MCF) problem is linear if  $QA$  is empty, and nonlinear (convex quadratic) if  $QA$  is nonempty.

The dual of the problem is:

$$\begin{aligned} \max \quad & \sum_{i \in N} P[i]b[i] - \sum_{(i,j) \in A} W[i, j]U[i, j] - \sum_{(i,j) \in AQ} V[i, j]^2/(2 * Q[i, j]) \\ (3.a) \quad & C[i, j] - P[i] + P[j] + W[i, j] - Z[i, j] = 0 \quad (i, j) \in AL \\ (3.b) \quad & C[i, j] - P[i] + P[j] + W[i, j] - Z[i, j] = V[i, j] \quad (i, j) \in AQ \\ (4.a) \quad & W[i, j] \geq 0 \quad (i, j) \in A \\ (4.b) \quad & Z[i, j] \geq 0 \quad (i, j) \in A \end{aligned}$$

$P[i]$  is said the vector of node potentials for the problem,  $W[i, j]$  are bound variables and  $Z[i, j]$  are slack variables. Given  $P[i]$ , the quantities

$$RC[i, j] = C[i, j] + Q[i, j] * X[i, j] - P[i] + P[j]$$

are said the "reduced costs" of arcs.

A primal and dual feasible solution pair is optimal if and only if the complementary slackness conditions

$$RC[i, j] > 0 \Rightarrow X[i, j] = 0$$

$$RC[i, j] < 0 \Rightarrow X[i, j] = U[i, j]$$

are satisfied for all arcs (i, j) of A.

The [MCFClass](#) class provides an interface with methods for managing and solving problems of this kind. Actually, the class can also be used as an interface for more general NonLinear MCF problems, where the cost function either nonseparable (  $C(X)$  ) or arc-separable (  $\sum_{(i,j) \in A} C_{i,j}(X[i, j])$  ). However, solvers for NonLinear MCF problems are typically objective-function-specific, and there is no standard way for inputting a nonlinear function different from a separable convex quadratic one, so only the simplest form is dealt with in the interface, leaving more complex NonLinear parts to the interface of derived classes.

## 7.3.2 Member Typedef Documentation

### 7.3.2.1 typedef double MCFClass::FONumber

type of the objective function: has to hold sums of products of FNumber(s) by CNumber(s)

## 7.3.3 Member Enumeration Documentation

### 7.3.3.1 enum MCFClass::MCFAnswer

Public enum describing the possible reoptimization status of the MCF solver.

**Enumerator:**

*kNo* no  
*kYes* yes

### 7.3.3.2 enum MCFClass::MCFFIFrmt

Public enum describing the possible file formats in [WriteMCF\(\)](#).

**Enumerator:**

*kDimacs* DIMACS file format for MCF.  
*kQDimacs* quadratic DIMACS file format for MCF  
*kMPS* MPS file format for LP.  
*kFWMPS* "Fixed Width" MPS format

### 7.3.3.3 enum MCFClass::MCFParam

Public enum describing the possible parameters of the MCF solver, to be used with the methods [SetPar\(\)](#) and [GetPar\(\)](#).

**Enumerator:**

*kMaxTime* max time

***kMaxIter*** max number of iteration

***kEpsFlw*** tolerance for flows

***kEpsDfct*** tolerance for deficits

***kEpsCst*** tolerance for costs

***kReopt*** whether or not to reoptimize

***kLastParam*** dummy parameter: this is used to allow derived classes to "extend" the set of parameters.

#### 7.3.3.4 enum MCFClass::MCFStatus

Public enum describing the possible status of the MCF solver.

**Enumerator:**

***kUnsolved*** no solution available

***kOK*** optimal solution found

***kStopped*** optimization stopped

***kUnfeasible*** problem is unfeasible

***kUnbounded*** problem is unbounded

***kError*** error in the solver

### 7.3.4 Constructor & Destructor Documentation

#### 7.3.4.1 MCFClass::MCFClass (cIndex *nm*x = 0, cIndex *mm*x = 0) [inline]

Constructor of the class.

*nm*x and *mm*x, if provided, are taken to be respectively the maximum number of nodes and arcs in the network. If nonzero values are passed, memory allocation can be anticipated in the constructor, which is sometimes desirable. The maximum values are stored in the protected fields *n*max and *m*max, and can be changed with [LoadNet\(\)](#) [see below]; however, changing them typically requires memory allocation/deallocation, which is sometimes undesirable outside the constructor.

After that an object has been constructed, no problem is loaded; this has to be done with [LoadNet\(\)](#) [see below]. Thus, it is an error to invoke any method which requires the presence of a problem (typicall all except those in the initializations part). The base class provides two protected fields *n* and *m* for the current number of nodes and arcs, respectively, that are set to 0 in the constructor precisely to indicate that no instance is currently loaded.

#### 7.3.4.2 virtual MCFClass::~~MCFClass () [inline, virtual]

Destructor of the class. The implementation in the base class only deletes the MCFt field. It is virtual, as it should be.

### 7.3.5 Member Function Documentation

#### 7.3.5.1 **virtual Index MCFClass::AddArc (cIndex *Start*, cIndex *End*, cFNumber *aU*, cCNumber *aC*) [pure virtual]**

Add the new arc ( *Start* , *End* ) with cost *aC* and capacity *aU*, returning its name. Inf<Index>() is returned if there is no room for a new arc. Remember that arc names go from 0 to mmax - 1.

Implemented in [MCFSimplex](#).

#### 7.3.5.2 **virtual Index MCFClass::AddNode (cFNumber *aDfct*) [pure virtual]**

Add a new node with deficit *aDfct*, returning its name. Inf<Index>() is returned if there is no room for a new node. Remember that the node names are either { 0 .. nmax - 1 } or { 1 .. nmax }, depending on the value of USENAME0.

Implemented in [MCFSimplex](#).

#### 7.3.5.3 **virtual void MCFClass::ChangeArc (cIndex *name*, cIndex *nSN* = Inf< Index > (), cIndex *nEN* = Inf< Index > ()) [pure virtual]**

Change the starting and/or ending node of arc '*name*' to *nSN* and *nEN*. Each parameter being Inf<Index>() means to leave the previous starting or ending node untouched. When this method is called '*name*' can be either the name of a "normal" arc or that of a "closed" arc [see [CloseArc\(\)](#) above]: in the latter case, at the end of [ChangeArc\(\)](#) the arc is \*still closed\*, and it remains so until [OpenArc\( name \)](#) [see above] is called.

Implemented in [MCFSimplex](#).

#### 7.3.5.4 **void MCFClass::CheckDSol (void) [inline]**

Check that the dual solution returned by the solver is dual feasible. (to within the tolerances set by SetPar(kEps\*\*\*\*) [see above], if any). Also, check that the objective function value is correct.

This method is implemented by the base class, using the above methods for collecting the solutions and the methods of the next section for reading the data of the problem; as such, they will work for any derived class that properly implements all these methods.

#### 7.3.5.5 **void MCFClass::CheckPSol (void) [inline]**

Check that the primal solution returned by the solver is primal feasible. (to within the tolerances set by SetPar(kEps\*\*\*\*) [see above], if any). Also, check that the objective function value is correct.

This method is implemented by the base class, using the above methods for collecting the solutions and the methods of the next section for reading the data of the problem; as such, they will work for any derived class that properly implements all these methods.

#### 7.3.5.6 **virtual void MCFClass::ChgCost (Index *arc*, cCNumber *NCost*) [pure virtual]**

Change the cost of the *i*-th arc.

#### Note

changing the costs of arcs that \*do not exist\* is \*not allowed\*; only arcs which have not been



closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

#### 7.3.5.7 **virtual void MCFClass::ChgCosts (cCRow *NCost*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Change the arc costs. In particular, change the costs that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- *\*and\** whose name belongs to the interval [‘strt’, ‘stp’).

That is, if  $strt \leq nms[i] < stp$ , then the  $nms[i]$ -th cost will be changed to  $NCost[i]$ . If  $nms == NULL$  (as the default), *\*all\** the entries in the given range will be changed; if  $stp > MCFm()$ , then the smaller bound is used.

#### Note

changing the costs of arcs that *\*do not exist\** is *\*not allowed\**; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

#### 7.3.5.8 **virtual void MCFClass::ChgDfct (Index *node*, cFNumber *NDfct*) [pure virtual]**

Change the deficit of the  $i$ -th node.

Note that, in [ChgDfct\[s\]\(\)](#), node "names" ( $i$ ,  $strt$ /  $stp$  or those contained in  $nms[]$ ) go from 0 to  $n - 1$ , regardless to the value of  $USERNAME0$ ; hence, if  $USERNAME0 == 0$  then the first node is "named 1", but its deficit can be changed e.g. with [ChgDfcts](#)( 0 ,  $new\_deficit$  ).

#### Note

changing the capacities of nodes that *\*do not exist\** is *\*not allowed\**; only nodes that have not been deleted [see [DelNode\(\)](#) below] can be touched with these methods.

#### 7.3.5.9 **virtual void MCFClass::ChgDfcts (cFRow *NDfct*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Change the node deficits. In particular, change the deficits that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- *\*and\** whose name belongs to the interval [‘strt’, ‘stp’).

That is, if  $strt \leq nms[i] < stp$ , then the  $nms[i]$ -th deficit will be changed to  $NDfct[i]$ . If  $nms == NULL$  (as the default), *\*all\** the entries in the given range will be changed; if  $stp > MCFn()$ , then the smaller bound is used.

Note that, in [ChgDfcts\(\)](#), node "names" ( $strt$ ,  $stp$  or those contained in  $nms[]$ ) go from 0 to  $n - 1$ , regardless to the value of  $USERNAME0$ ; hence, if  $USERNAME0 == 0$  then the first node is "named 1", but its deficit can be changed e.g. with [ChgDfcts](#)(  $\&new\_deficit$  ,  $NULL$  , 0 , 1 ).

**Note**

changing the capacities of nodes that *\*do not exist\** is *\*not allowed\**; only nodes that have not been deleted [see [DelNode\(\)](#) below] can be touched with these methods.

### 7.3.5.10 **virtual void MCFClass::ChgQCoef (Index *arc*, cCNumber *NQCoef*) [inline, virtual]**

**Parameters**

*NQCoef* Change the quadratic coefficient of the cost of the *i*-th arc.

Note that the method is *\*not\** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

**Note**

changing the costs of arcs that *\*do not exist\** is *\*not allowed\**; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

### 7.3.5.11 **virtual void MCFClass::ChgQCoef (cCRow *NQCoef* = NULL, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf<Index>()) [inline, virtual]**

**Parameters**

*stp* Change the quadratic coefficients of the arc costs. In particular, change the coefficients that are:

- listed in into the vector of indices '*nms*' (ordered in increasing sense and Inf<Index>()-terminated),
- *\*and\** whose name belongs to the interval [*'strt'*, *'stp'*).

That is, if *strt* ≤ *nms*[ *i* ] < *stp*, then the *nms*[ *i* ]-th cost will be changed to *NCost*[ *i* ]. If *nms* == NULL (as the default), *\*all\** the entries in the given range will be changed; if *stp* > [MCFm\(\)](#), then the smaller bound is used. If *NQCoef* == NULL, all the specified coefficients are set to zero.

Note that the method is *\*not\** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

**Note**

changing the costs of arcs that *\*do not exist\** is *\*not allowed\**; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

### 7.3.5.12 **virtual void MCFClass::ChgUCap (Index *arc*, cFNumber *NCap*) [pure virtual]**

Change the capacity of the *i*-th arc.

**Note**

changing the capacities of arcs that *\*do not exist\** is *\*not allowed\**; only arcs that have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

### 7.3.5.13 virtual void MCFClass::ChgUCaps (cFRow *NCap*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]

Change the arc capacities. In particular, change the capacities that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- \*and\* whose name belongs to the interval [*strt*, *stp*).

That is, if  $strt \leq nms[i] < stp$ , then the  $nms[i]$ -th capacity will be changed to  $NCap[i]$ . If  $nms == NULL$  (as the default), \*all\* the entries in the given range will be changed; if  $stp > MCFm()$ , then the smaller bound is used.

#### Note

changing the capacities of arcs that \*do not exist\* is \*not allowed\*; only arcs that have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C\_INF costs] can be touched with these methods.

### 7.3.5.14 virtual void MCFClass::CloseArc (cIndex *name*) [pure virtual]

"Close" the arc ‘name’. Although all the associated information (name, cost, capacity, end and start node) is kept, the arc is removed from the problem until [OpenArc\(i\)](#) [see below] is called.

"closed" arcs always have 0 flow, but are otherwise counted as any other arc; for instance, [MCFm\(\)](#) does \*not\* decrease as an effect of a call to [CloseArc\(\)](#). How this closure is implemented is solver-specific.

Implemented in [MCFSimplex](#).

### 7.3.5.15 virtual void MCFClass::DelArc (cIndex *name*) [pure virtual]

Delete the arc ‘name’. Unlike "closed" arcs, all the information associated with a deleted arc is lost and ‘name’ is made available as a name for new arcs to be created with [AddArc\(\)](#) [see below].

It furthermore ‘name’ is the last arc, the number of arcs as reported by [MCFm\(\)](#) is reduced by at least one, until the m-th arc is not a deleted one. Otherwise, the flow on the arc is always ensured to be 0.

Implemented in [MCFSimplex](#).

### 7.3.5.16 virtual void MCFClass::DelNode (cIndex *name*) [pure virtual]

Delete the node ‘name’.

For any value of ‘name’, all incident arcs to that node are closed [see [CloseArc\(\)](#) above] (\*not\* Deleted, see [DelArc\(\)](#) below) and the deficit is set to zero.

It furthermore ‘name’ is the last node, the number of nodes as reported by [MCFn\(\)](#) is reduced by at least one, until the n-th node is not a deleted one.

Implemented in [MCFSimplex](#).

### 7.3.5.17 template<class T> bool MCFClass::ETZ (T *x*, const T *eps*) [inline, protected]

true if flow *x* is equal to zero (possibly considering tolerances).

**7.3.5.18 void MCFClass::GetPar (int *par*, double & *val*) [inline, virtual]**

This method returns one of the integer parameter of the algorithm.

**Parameters**

*par* is the parameter to return [see SetPar( double ) for comments];

*val* upon return, it will contain the value of the parameter.

The base class implementation handles the parameters kEpsFlw, kEpsDfct, kEpsCst, and kMaxTime.

**7.3.5.19 void MCFClass::GetPar (int *par*, int & *val*) [inline, virtual]**

This method returns one of the integer parameter of the algorithm.

**Parameters**

*par* is the parameter to return [see SetPar( int ) for comments];

*val* upon return, it will contain the value of the parameter.

The base class implementation handles the parameters kMaxIter and kReopt.

**7.3.5.20 template<class T > bool MCFClass::GEZ (T *x*, const T *eps*) [inline, protected]**

true if flow *x* is greater than or equal to zero (possibly considering tolerances).

**7.3.5.21 template<class T > bool MCFClass::GT (T *x*, T *y*, const T *eps*) [inline, protected]**

true if flow *x* is greater than flow *y* (possibly considering tolerances).

**7.3.5.22 template<class T > bool MCFClass::GTZ (T *x*, const T *eps*) [inline, protected]**

true if flow *x* is greater than zero (possibly considering tolerances).

**7.3.5.23 virtual bool MCFClass::HaveNewPi (void) [inline, virtual]**

Return true if a different (approximately) optimal dual solution is available. If the method returns true, then any subsequent call to (any form of) [MCFGetPi\(\)](#) will return a different dual solution, and [MCFGetRC\(\)](#) [see below] will return the corresponding reduced costs. The new solution need not be optimal (although, ideally, it has to be "good"); this can be checked by comparing its objective function value, that will be returned by a call to [MCFGetDFO\(\)](#) [see below].

Any subsequent call of [HaveNewPi\(\)](#) that returns true produces a new solution, until the first that returns false; from then on, no new solutions will be generated until something changes in the problem's data.

Note that a default implementation of [HaveNewPi\(\)](#) is provided which is good for those solvers that only produce one optimal dual solution.

**7.3.5.24 virtual bool MCFClass::HaveNewX(void) [inline, virtual]**

Return true if a different (approximately) optimal primal solution is available. If the method returns true, then any subsequent call to (any form of) [MCFGetX\(\)](#) will return a different primal solution w.r.t. the one that was being returned \*before\* the call to [HaveNewX\(\)](#). This solution need not be optimal (although, ideally, it has to be "good"); this can be checked by comparing its objective function value, that will be returned by a call to [MCFGetFO\(\)](#) [see below].

Any subsequent call of [HaveNewX\(\)](#) that returns true produces a new solution, until the first that returns false; from then on, no new solutions will be generated until something changes in the problem's data.

Note that a default implementation of [HaveNewX\(\)](#) is provided which is good for those solvers that only produce one optimal primal solution.

**7.3.5.25 virtual bool MCFClass::IsClosedArc(cIndex name) [pure virtual]**

[IsClosedArc\(\)](#) returns true if and only if the arc 'name' is closed.

Implemented in [MCFSimplex](#).

**7.3.5.26 virtual bool MCFClass::IsDeletedArc(cIndex name) [pure virtual]**

Return true if and only if the arc 'name' is deleted. It should only be called with name < [MCFm\(\)](#), as every other arc is deleted by definition.

Implemented in [MCFSimplex](#).

**7.3.5.27 template<class T> bool MCFClass::LEZ(T x, const T eps) [inline, protected]**

true if flow x is less than or equal to zero (possibly considering tolerances).

**7.3.5.28 void MCFClass::LoadDMX(istream & DMXs, bool IsQuad = false) [inline, virtual]**

Read a MCF instance in DIMACS standard format from the istream. The format is the following. The first line must be

```
p min <number of="" nodes>=""> <number of="" arcs>="">
```

Then the node definition lines must be found, in the form

```
n <node number>=""> <node supply>="">
```

Not all nodes need have a node definition line; these are given zero supply, i.e., they are transshipment nodes (supplies are the inverse of deficits, i.e., a node with positive supply is a source node). Finally, the arc definition lines must be found, in the form

```
a <start node>=""> <end node>=""> <lower bound>=""> <upper bound>=""> <flow cost>="">
```

There must be exactly <number of="" arcs>=""> arc definition lines in the file.

This method is \*not\* pure virtual because an implementation is provided by the base class, using the [LoadNet\(\)](#) method (which \*is\* pure virtual). However, the method \*is\* virtual to allow derived classes to implement more efficient versions, should they have any reason to do so.

**Note**

Actually, the file format accepted by LoadDMX (at least in the base class implementation) is more general than the DIMACS standard format, in that it is allowed to mix node and arc definitions in any order, while the DIMACS file requires all node information to appear before all arc information.

Other than for the above, this method is assumed to allow for quadratic\* Dimacs files, encoding for convex quadratic separable Min Cost Flow instances. This is a simple extension where each arc descriptor has a sixth field, `<quadratic cost>=""`. The provided istream is assumed to be quadratic Dimacs file if `IsQuad` is true, and a regular linear Dimacs file otherwise.

**7.3.5.29** `virtual void MCFClass::LoadNet (cIndex nmx = 0, cIndex mmx = 0, cIndex pn = 0, cIndex pm = 0, cFRow pU = NULL, cCRow pC = NULL, cFRow pDfct = NULL, cIndex_Set pSn = NULL, cIndex_Set pEn = NULL) [pure virtual]`

Inputs a new network.

The parameters `nmx` and `mmx` are the new max number of nodes and arcs, possibly overriding those set in the constructor [see above], although at the likely cost of memory allocation and deallocation. Passing `nmx == mmx == 0` is intended as a signal to the solver to deallocate everything and wait for new orders; in this case, all the other parameters are ignored.

Otherwise, in principle all the other parameters have to be provided. Actually, some of them may not be needed for special classes of MCF problems (e.g., costs in a MaxFlow problem, or start/end nodes in a problem defined over a graph with fixed topology, such as a complete graph). Also, passing `NULL` is allowed to set default values.

The meaning of the parameters is the following:

- `pn` is the current number of nodes of the network ( $\leq$  `nmax`).
- `pm` is the number of arcs of the network ( $\leq$  `mmax`).
- `pU` is the `m`-vector of the arc upper capacities; capacities must be nonnegative, but can in principle be infinite ( $==$  `F_INF`); passing `pU == NULL` means that all capacities are infinite;
- `pC` is the `m`-vector of the arc costs; costs must be finite ( $<$  `C_INF`); passing `pC == NULL` means that all costs must be 0.
- `pDfct` is the `n`-vector of the node deficits; source nodes have negative deficits and sink nodes have positive deficits; passing `pDfct == NULL` means that all deficits must be 0 (a circulation problem);
- `pSn` is the `m`-vector of the arc starting nodes; `pSn == NULL` is in principle not allowed, unless the topology of the graph is fixed;
- `pEn` is the `m`-vector of the arc ending nodes; same comments as for `pSn`.

Note that node "names" in the arrays `pSn` and `pEn` must go from 1 to `pn` if the macro `USANAME0` [see above] is set to 0, while they must go from 0 to `pn - 1` if `USANAME0` is set to 1. In both cases, however, the deficit of the first node is read from the first (0-th) position of `pDfct`, that is if `USANAME0 == 0` then the deficit of the node with name 'i' is read from `pDfct[i - 1]`.

The data passed to `LoadNet()` can be used to specify that the arc 'i' must not "exist" in the problem. This is done by passing `pC[i] == C_INF`; solvers which don't read costs are forced to read them in order to check this, unless they provide alternative solver-specific ways to accomplish the same tasks. These arcs

are "closed", as for the effect of [CloseArc\(\)](#) [see below]. "invalid" costs ( $\text{C\_INF}$ ) are set to 0 in order to being subsequently capable of "opening" them back with [OpenArc\(\)](#) [see below]. The way in which these non-existent arcs are physically dealt with is solver-specific; in some solvers, for instance, this could be obtained by simply putting their capacity to zero. Details about these issues should be found in the interface of derived classes.

Note that the quadratic part of the objective function, if any, is not dealt with in [LoadNet\(\)](#); it can only be separately provided with [ChgQCoef\(\)](#) [see below]. By default, the problem is linear, i.e., all coefficients of the second-order terms in the objective function are assumed to be zero.

Implemented in [MCFSimplex](#).

**7.3.5.30** `template<class T> bool MCFClass::LT (T x, T y, const T eps) [inline, protected]`

true if flow x is less than flow y (possibly considering tolerances).

**7.3.5.31** `template<class T> bool MCFClass::LTZ (T x, const T eps) [inline, protected]`

true if flow x is less than zero (possibly considering tolerances).

**7.3.5.32** `virtual void MCFClass::MCFArcs (Index_Set Startv, Index_Set Endv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]`

Write the starting (tail) and ending (head) nodes of the arcs in Startv[] and Endv[]. If nms == NULL, then the information relative to all arcs is written into Startv[] and Endv[], otherwise Startv[ i ] and Endv[ i ] contain the information relative to arc nms[ i ] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with  $\text{strt} \leq i < \min(\text{MCFm}(), \text{stp})$ . 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to arcs which are \*both\* in nms and whose index is in the correct range are returned.

Startv or Endv can be NULL, meaning that only the other information is required.

#### Note

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

If the graph is "dynamic", be careful to use [MCFn\(\)](#) e [MCFm\(\)](#) to properly choose the dimension of nodes and arcs arrays.

**7.3.5.33** `virtual CNumber MCFClass::MCFCost (cIndex i) [pure virtual]`

Return the cost of the i-th arc.

Implemented in [MCFSimplex](#).

**7.3.5.34** `virtual cCRow MCFClass::MCFCosts (void) [inline, virtual]`

Return a read-only pointer to an internal vector containing the arc costs. Since this may \*not always be available\*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

### 7.3.5.35 **virtual void MCFClass::MCFCosts (CRow *Costv*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Write the arc costs into *Costv*[]. If *nms* == NULL, then all the costs are written, otherwise *Costv*[ *i* ] contains the information relative to arc *nms*[ *i* ] (*nms*[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with *strt* <= *i* < min( [MCFm\(\)](#) , *stp* ). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to arcs which are \*both\* in *nms* and whose index is in the correct range are returned.

### 7.3.5.36 **virtual FNumber MCFClass::MCFDfct (cIndex *i*) [pure virtual]**

Return the deficit of the *i*-th node.

#### Note

Here node "names" go from 0 to *n* - 1, regardless to the value of *USERNAME0*; hence, if *USERNAME0* == 0 then the first node is "named 1", but its deficit is returned by [MCFDfct\( 0 \)](#).

Implemented in [MCFSimplex](#).

### 7.3.5.37 **virtual cFRow MCFClass::MCFDfcts (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the node deficits. Since this may \*not\* always be available\*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

#### Note

Here node "names" go from 0 to *n* - 1, regardless to the value of *USERNAME0*; hence, if *USERNAME0* == 0 then the first node is "named 1", but its deficit is contained in [MCFDfcts\(\)\[ 0 \]](#).

### 7.3.5.38 **virtual void MCFClass::MCFDfcts (FRow *Dfctv*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Write the node deficits into *Dfctv*[]. If *nms* == NULL, then all the deficits are written, otherwise *Dfctv*[ *i* ] contains the information relative to node *nms*[ *i* ] (*nms*[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the nodes 'i' with *strt* <= *i* < min( [MCFn\(\)](#) , *stp* ). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to nodes which are \*both\* in *nms* and whose index is in the correct range are returned.

#### Note

Here node "names" (*strt* and *stp*, those contained in *nms*[] or 'i' in [MCFDfct\(\)](#)) go from 0 to *n* - 1, regardless to the value of *USERNAME0*; hence, if *USERNAME0* == 0 then the first node is "named 1", but its deficit is returned by [MCFDfcts\( \*Dfctv\* , NULL , 0 , 1 \)](#).



**7.3.5.39 virtual Index MCFClass::MCFENde (cIndex i) [pure virtual]**

Return the ending (head) node of the arc 'i'.

**Note**

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

Implemented in [MCFSimplex](#).

**7.3.5.40 virtual cIndex\_Set MCFClass::MCFENdes (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the ending (head) nodes for each arc. Since this may \*not always be available\*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

**7.3.5.41 virtual FNumber MCFClass::MCFGetDFO (void) [inline, virtual]**

Return the objective function value of the dual solution currently returned by [MCFGetPi\(\)](#) / [MCFGetRC\(\)](#). This value (possibly) changes after any call to [HaveNewPi\(\)](#) that returns true. The relations between [MCFGetStatus\(\)](#) and [MCFGetDFO\(\)](#) are analogous to these of [MCFGetFO\(\)](#), except that a finite value corresponding to kStopped must be a lower bound on the optimal objective function value (typically, the objective function value one dual feasible solution).

A default implementation is provided for [MCFGetDFO\(\)](#), which is good for MCF solvers where the primal and dual optimal solution values always are identical (except if the problem is unfeasible/unbounded).

**7.3.5.42 virtual FNumber MCFClass::MCFGetFO (void) [pure virtual]**

Return the objective function value of the primal solution currently returned by [MCFGetX\(\)](#).

If [MCFGetStatus\(\)](#) == kOK, this is guaranteed to be the optimal objective function value of the problem (to within the optimality tolerances), but only prior to any call to [HaveNewX\(\)](#) that returns true. [MCFGetFO\(\)](#) typically returns Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnfeasible and

- Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnbounded. If [MCFGetStatus\(\)](#) == kStopped and [MCFGetFO\(\)](#) returns a finite value, it must be an upper bound on the optimal objective function value (typically, the objective function value of one primal feasible solution).

Implemented in [MCFSimplex](#).

**7.3.5.43 virtual cCRow MCFClass::MCFGetPi (void) [inline, virtual]**

Return a read-only pointer to an internal data structure containing the node potentials. Since this may \*not always be available\*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

#### 7.3.5.44 **virtual void MCFClass::MCFGetPi (CRow *P*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Writes the optimal node potentials in the vector *P*[]. If *nms* == NULL, the node potential of node 'i' (i in 0 .. n - 1) is written in *P*[ i ] (note that here node names always start from zero, regardless to the value of *USENAME0*). If *nms* != NULL, it must point to a vector of indices in 0 .. n - 1 (ordered in increasing sense and Inf<Index>()-terminated), and the node potential of *nms*[ i ] is written in *P*[ i ]. Note that, unlike [MCFGetX\(\)](#) above, *nms* is an *\*input\** of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the nodes 'i' with *strt* <= i < min( [MCFn\(\)](#) , *stp* ). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to nodes which are *\*both\** in *nms*[] and whose index is in the correct range are returned.

#### 7.3.5.45 **virtual CNumber MCFClass::MCFGetRC (cIndex *i*) [pure virtual]**

Return the reduced cost of the i-th arc. This information should be cheaply available in most implementations.

##### Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Implemented in [MCFSimplex](#).

#### 7.3.5.46 **virtual cCRow MCFClass::MCFGetRC (void) [inline, virtual]**

Return a read-only pointer to an internal data structure containing the reduced costs. Since this may *\*not\** always be available\*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

##### Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

#### 7.3.5.47 **virtual void MCFClass::MCFGetRC (CRow *CR*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Write the reduced costs corresponding to the current dual solution in *RC*[]. If *nms* == NULL, the reduced cost of arc 'i' (i in 0 .. m - 1) is written in *RC*[ i ]; if *nms* != NULL, it must point to a vector of indices in 0 .. m - 1 (ordered in increasing sense and Inf<Index>()-terminated), and the reduced cost of arc *nms*[ i ] is written in *RC*[ i ]. Note that, unlike [MCFGetX\(\)](#) above, *nms* is an *\*input\** of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with *strt* <= i < min( [MCFm\(\)](#) , *stp* ). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to arcs which are *\*both\** in *nms*[] and whose index is in the correct range are returned.

##### Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

**7.3.5.48 virtual MCFStatePtr MCFClass::MCFGetState (void) [inline, virtual]**

Save the state of the MCF solver. The [MCFClass](#) interface supports the notion of saving and restoring the state of the MCF solver, such as the current/optimal basis in a simplex solver. The "empty" class [MCFState](#) is defined as a placeholder for state descriptions.

[MCFGetState\(\)](#) creates and returns a pointer to an object of (a proper derived class of) class [MCFState](#) which describes the current state of the MCF solver.

**7.3.5.49 int MCFClass::MCFGetStatus (void) [inline]**

Returns an int describing the current status of the MCF solver. Possible return values are:

- `kUnSolved` [SolveMCF\(\)](#) has not been called yet, or the data of the problem has been changed since the last call;
- `kOK` optimization has been carried out successfully;
- `kStopped` optimization have been stopped before that the stopping conditions of the solver applied, e.g. because of the maximum allowed number of "iterations" [see [SetPar\( int \)](#)] or the maximum allowed time [see [SetPar\( double \)](#)] has been reached; this is not necessarily an error, as it might just be required to re-call [SolveMCF\(\)](#) giving it more "resources" in order to solve the problem;
- `kUnfeasible` if the current MCF instance is (primal) unfeasible;
- `kUnbounded` if the current MCF instance is (primal) unbounded (this can only happen if the solver actually allows `F_INF` capacities, which is nonstandard in the interface);
- `kError` if there was an error during the optimization; this typically indicates that computation cannot be resumed, although solver-dependent ways of dealing with solver-dependent errors may exist.

[MCFClass](#) has a protected `int` member `status` that can be used by derived classes to hold status information and that is returned by the standard implementation of this method. Note that `status` is an `int` and not an `enum`, and that an `int` is returned by this method, in order to allow the derived classes to extend the set of return values if they need to do so.

**7.3.5.50 virtual Index MCFClass::MCFGetUnbCycl (Index\_Set Pred, Index\_Set ArcPred) [inline, virtual]**

Return an unboundedness certificate. In an unbounded MCF problem, unboundedness can always be reduced to the existence of a directed cycle with negative cost and all arcs having infinite capacity. When detecting unboundedness, MCF solvers are typically capable to provide one such cycle. This information can be useful, and this method is provided for getting it. It can be called only if [MCFGetStatus\(\)](#) == `kUnbounded`, and writes in `Pred[]` and `ArcPred[]`, respectively, the node and arc predecessor function of the cycle. That is, if node 'i' belongs to the cycle then '`Pred[ i ]`' contains the name of the predecessor of 'j' of 'i' in the cycle (note that node names depend on `USERNAME0`), and '`ArcPred[ i ]`' contains the index of the arc joining the two (note that in general there may be multiple copies of each arc). Entries of the vectors for nodes not belonging to the cycle are in principle undefined, and the name of one node belonging to the cycle is returned by the method. Note that if there are multiple cycles with negative costs this method will return just one of them (finding the cycle with most negative cost is an NO-hard problem), although solvers should be able to produce cycles with "large negative" cost.

However, not all solvers may be (easily) capable of providing this information; thus, returning `Inf<Index>()` is allowed, as in the base class implementation, to signify that this information is not available.

#### 7.3.5.51 `virtual FNumber MCFClass::MCFGetUnfCut (Index_Set Cut) [inline, virtual]`

Return an unfeasibility certificate. In an unfeasible MCF problem, unfeasibility can always be reduced to the existence of a cut (subset of nodes of the graph) such as either:

- the inverse of the deficit of the cut (the sum of all the deficits of the nodes in the cut) is larger than the forward capacity of the cut (sum of the capacities of forward arcs in the cut); that is, the nodes in the cut globally produce more flow than can be routed to sinks outside the cut;
- the deficit of the cut is larger than the backward capacity of the cut (sum of the capacities of backward arcs in the cut); that is, the nodes in the cut globally require more flow than can be routed to them from sources outside the cut.

When detecting unfeasibility, MCF solvers are typically capable to provide one such cut. This information can be useful - typically, the only way to make the problem feasible is to increase the capacity of at least one of the forward/backward arcs of the cut -, and this method is provided for getting it. It can be called only if `MCFGetStatus() == kUnfeasible`, and should write in `Cut` the set of names of nodes in the unfeasible cut (note that node names depend on `USENAME0`), `Inf<Index>()`-terminated, returning the deficit of the cut (which allows to distinguish which of the two cases above hold). In general, no special properties can be expected from the returned cut, but solvers should be able to provide e.g. "small" cuts.

However, not all solvers may be (easily) capable of providing this information; thus, returning 0 (no cut) is allowed, as in the base class implementation, to signify that this information is not available.

#### 7.3.5.52 `virtual cFRow MCFClass::MCFGetX (void) [inline, virtual]`

Return a read-only pointer to an internal data structure containing the flow solution in "dense" format. Since this may *\*not* always be available\*, depending on the implementation, this method can (uniformly) return `NULL`. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

#### 7.3.5.53 `virtual void MCFClass::MCFGetX (FRow F, Index_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]`

Write the optimal flow solution in the vector `F[]`. If `nms == NULL`, `F[]` will be in "dense" format, i.e., the flow relative to arc 'i' (i in 0 .. m - 1) is written in `F[ i ]`. If `nms != NULL`, `F[]` will be in "sparse" format, i.e., the indices of the nonzero elements in the flow solution are written in `nms` (that is then `Inf<Index>()`-terminated) and the flow value of arc `nms[ i ]` is written in `F[ i ]`. Note that `nms` is not\* guaranteed to be ordered. Also, note that, unlike `MCFGetRC()` and `MCFGetPi()` [see below], `nms` is an *\*output\** of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with `strt <= i < min( MCFm(), stp )`.

#### 7.3.5.54 `Index MCFClass::MCFm (void) [inline]`

Return the number of arcs in the current graph. The implementation of the method in the base class returns the protected fields `m`, which is provided for derived classes to hold this information.

**7.3.5.55 Index MCFClass::MCFmmax (void) [inline]**

Return the maximum number of arcs for this instance of [MCFClass](#). The implementation of the method in the base class returns the protected fields `mmax`, which is provided for derived classes to hold this information.

**7.3.5.56 Index MCFClass::MCFn (void) [inline]**

Return the number of nodes in the current graph. The implementation of the method in the base class returns the protected fields `n`, which is provided for derived classes to hold this information.

**7.3.5.57 Index MCFClass::MCFnmax (void) [inline]**

Return the maximum number of nodes for this instance of [MCFClass](#). The implementation of the method in the base class returns the protected fields `nmax`, which is provided for derived classes to hold this information.

**7.3.5.58 virtual void MCFClass::MCFPutState (MCFStatePtr S) [inline, virtual]**

Restore the solver to the state in which it was when the state 'S' was created with [MCFGetState\(\)](#) [see above].

The typical use of this method is the following: a MCF problem is solved and the "optimal state" is set aside. Then the problem changes and it is re-solved. Then, the problem has to be changed again to a form that is close to the original one but rather different from the second one (think of a long backtracking in a Branch & Bound) to which the current "state" refers. Then, the old optimal state can be expected to provide a better starting point for reoptimization [see [ReOptimize\(\)](#) below].

Note, however, that the state is only relative to the optimization process, i.e., this operation is meaningless if the data of the problem has changed in the meantime. So, if a state has to be used for speeding up reoptimization, the following has to be done:

- first, the data of the solver is brought back to *\*exactly\** the same as it was at the moment where the state 'S' was created (prior than this operation a `ReOptimize( false )` call is probably advisable);
- then, [MCFPutState\(\)](#) is called (and `ReOptimize( true )` is called);
- only afterwards the data of the problem is changed to the final state and the problem is solved.

A "put state" operation does not "deplete" the state, which can therefore be used more than once. Indeed, a state is constructed inside the solver for each call to [MCFGetState\(\)](#), but the solver never deletes statuses; this has to be done on the outside when they are no longer needed (the solver must be "resistent" to deletion of the state at any moment).

Since not all the MCF solvers reoptimize (efficiently enough to make these operations worth), an "empty" implementation that does nothing is provided by the base class.

**7.3.5.59 virtual cCRow MCFClass::MCFQCoef (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the arc costs. Since this may *\*not always be available\**, depending on the implementation, this method can (uniformly) return `NULL`. This is done by the base class already, so a derived class that does not have the information ready (such as "pure linear" MCF solvers that only work with all zero quadratic coefficients) does not need to implement the method.

### 7.3.5.60 **virtual CNumber MCFClass::MCFQCoef (cIndex *i*) [inline, virtual]**

#### Parameters

*i* Return the quadratic coefficients of the cost of the *i*-th arc. Note that the method is *\*not\** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

### 7.3.5.61 **virtual void MCFClass::MCFQCoef (CRow *Qv*, cIndex\_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf<Index>()) [inline, virtual]**

#### Parameters

*stp* Write the quadratic coefficients of the arc costs into *Qv*[]. If *nms* == NULL, then all the costs are written, otherwise *Costv*[ *i* ] contains the information relative to arc *nms*[ *i* ] (*nms*[] must be Inf<Index>()-terminated).

The parameters '*strt*' and '*stp*' allow to restrict the output of the method to all and only the arcs '*i*' with  $strt \leq i < \min(\text{MCFm}(), stp)$ . '*strt*' and '*stp*' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to arcs which are *\*both\** in *nms* and whose index is in the correct range are returned.

Note that the method is *\*not\** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

### 7.3.5.62 **virtual Index MCFClass::MCFSNde (cIndex *i*) [pure virtual]**

Return the starting (tail) node of the arc '*i*'.

#### Note

If *USENAME0* == 0 then the returned node names will be in the range 1 .. *n*, while if *USENAME0* == 1 the returned node names will be in the range 0 .. *n* - 1.

Implemented in [MCFSimplex](#).

### 7.3.5.63 **virtual cIndex\_Set MCFClass::MCFSNdes (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the starting (tail) nodes for each arc. Since this may *\*not always be available\**, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

### 7.3.5.64 **virtual FNumber MCFClass::MCFUCap (cIndex *i*) [pure virtual]**

Return the capacity of the *i*-th arc.

Implemented in [MCFSimplex](#).

### 7.3.5.65 **virtual cFRow MCFClass::MCFUCaps (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the arc capacities. Since this may *\*not always be available\**, depending on the implementation, this method can (uniformly) return NULL. This is done

by the base class already, so a derived class that does not have the information ready does not need to implement the method.

**7.3.5.66** `virtual void MCFClass::MCFUCaps (FRow UCapv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]`

Write the arc capacities into UCapv[]. If nms == NULL, then all the capacities are written, otherwise UCapv[ i ] contains the information relative to arc nms[ i ] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with  $strt \leq i < \min(\text{MCFm}(), \text{stp})$ . 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to arcs which are \*both\* in nms and whose index is in the correct range are returned.

**7.3.5.67** `virtual void MCFClass::OpenArc (cIndex name) [pure virtual]`

Restore the previously closed arc 'name'. It is an error to open an arc that has not been previously closed. Implemented in [MCFSimplex](#).

**7.3.5.68** `virtual void MCFClass::PreProcess (void) [inline, virtual]`

Extract a smaller/easier equivalent MCF problem. The data of the instance is changed and the easier one is solved instead of the original one. In the MCF case, preprocessing may involve reducing bounds, identifying disconnected components of the graph etc. However, preprocessing is solver-specific.

This method can be implemented by derived classes in their solver-specific way. Preprocessing may reveal unboundedness or unfeasibility of the problem; if that happens, [PreProcess\(\)](#) should properly set the 'status' field, that can then be read with [MCFGetStatus\(\)](#) [see below].

Note that preprocessing may destroy all the solution information. Also, it may be allowed to change the data of the problem, such as costs/capacities of the arcs.

A valid preprocessing is doing nothing, and that's what the default implementation of this method (that is \*not\* pure virtual) does.

**7.3.5.69** `virtual void MCFClass::SetMCFTIME (bool TimeIt = true) [inline, virtual]`

Allocate an OPTtimers object [see OPTtypes.h] to be used for timing the methods of the class. The time can be read with [TimeMCF\(\)](#) [see below]. By default, or if SetMCFTIME( false ) is called, no timing is done. Note that, since all the relevant methods of the class are pure virtual, [MCFClass](#) can only manage the OPTtimers object, but it is due to derived classes to actually implement the timing.

#### Note

time accumulates over the calls: calling [SetMCFTIME\(\)](#), however, resets the counters, allowing to time specific groups of calls.  
of course, setting kMaxTime [see [SetPar\(\)](#) above] to any nonzero value has no effect unless SetMCFTIME( true ) has been called.

**7.3.5.70** `void MCFClass::SetPar (int par, double val) [inline, virtual]`

Set float parameters of the algorithm.

**Parameters**

**par** is the parameter to be set; the enum `MCFParam` can be used, but 'par' is an int (every enum is an int) so that the method can be extended by derived classes for the setting of their parameters

**value** is the value to assign to the parameter.

The base class implementation handles these parameters:

- `kEpsFlw`: sets the tolerance for controlling if the flow on an arc is zero to val. This also sets the tolerance for controlling if a node deficit is zero (see `kEpsDfct`) to  $\text{val} * < \text{max number of nodes} >$ ; this value should be safe for graphs in which any node has less than  $< \text{max number of nodes} >$  adjacent nodes, i.e., for all graphs but for very dense ones with "parallel arcs"
- `kEpsDfct`: sets the tolerance for controlling if a node deficit is zero to val, in case a better value than that automatically set by `kEpsFlw` (see above) is available (e.g.,  $\text{val} * k$  would be good if no node has more than  $k$  neighbours)
- `kEpsCst`: sets the tolerance for controlling if the reduced cost of an arc is zero to val. A feasible solution satisfying eps-complementary slackness, i.e., such that

$$RC[i, j] < -\text{eps} \Rightarrow X[i, j] = U[i, j]$$

and

$$RC[i, j] > \text{eps} \Rightarrow X[i, j] == 0,$$

is known to be  $(\text{eps} * n)$ -optimal.

- `kMaxTime`: sets the max time (in seconds) in which the MCF Solver can find an optimal solution (default 0, which means no limit).

Reimplemented in [MCFSimplex](#).

**7.3.5.71 void MCFClass::SetPar (int par, int val) [inline, virtual]**

Set integer parameters of the algorithm.

**Parameters**

**par** is the parameter to be set; the enum `MCFParam` can be used, but 'par' is an int (every enum is an int) so that the method can be extended by derived classes for the setting of their parameters

**value** is the value to assign to the parameter.

The base class implementation handles these parameters:

- `kMaxIter`: the max number of iterations in which the MCF Solver can find an optimal solution (default 0, which means no limit)
- `kReopt`: tells the solver if it has to reoptimize. The implementation in the base class sets a flag, the protected `bool` field `Senstv`; if true (default) this field instructs the MCF solver to try to exploit the information about the latest optimal solution to speedup the optimization of the current problem, while if the field is false the MCF solver should restart the optimization "from scratch" discarding any previous information. Usually reoptimization speeds up the computation considerably, but this is not always true, especially if the data of the problem changes a lot.

Reimplemented in [MCFSimplex](#).



**7.3.5.72 virtual void MCFCClass::SolveMCF (void) [pure virtual]**

Solver of the Min Cost Flow Problem. Attempts to solve the MCF instance currently loaded in the object. Implemented in [MCFSimplex](#).

**7.3.5.73 double MCFCClass::TimeMCF (void) [inline]**

Like TimeMCF(double,double) [see above], but returns the total time.

**7.3.5.74 void MCFCClass::TimeMCF (double & *t\_us*, double & *t\_ss*) [inline]**

Time the code. If called within any of the methods of the class that are "actively timed" (this depends on the subclasses), this method returns the user and system time (in seconds) since the start of that method. If methods that are actively timed call other methods that are actively timed, [TimeMCF\(\)](#) returns the (...) time since the beginning of the outer\* actively timed method. If called outside of any actively timed method, this method returns the (...) time spent in all the previous executions of all the actively timed methods of the class.

Implementing the proper calls to `MCFT->Start()` and `MCFT->Stop()` is due to derived classes; these should at least be placed at the beginning and at the end, respectively, of [SolveMCF\(\)](#) and presumably the `Chg***()` methods, that is, at least these methods should be "actively timed".

**7.3.5.75 void MCFCClass::WriteMCF (ostream & *oStrm*, int *frmt* = 0) [inline, virtual]**

Write the current MCF problem to an ostream. This may be useful e.g. for debugging purposes.

The base [MCFCClass](#) class provides output in two different formats, depending on the value of the parameter `frmt`:

- `kDimacs` the problem is written in DIMACS standard format, read by most MCF codes available;
- `kMPS` the problem is written in the "modern version" (tab-separated) of the MPS format, read by most LP/MIP solvers;
- `kFWMPS` the problem is written in the "old version" (fixed width fields) of the MPS format; this is read by most LP/MIP solvers, but some codes still require the old format.

The implementation of [WriteMCF\(\)](#) in the base class uses all the above methods for reading the data; as such it will work for any derived class that properly implements this part of the interface, but it may not be very efficient. Thus, the method is virtual to allow the derived classes to either re-implement [WriteMCF\(\)](#) for the above two formats in a more efficient way, and/or to extend it to support other solver-specific formats.

**Note**

None of the above two formats supports quadratic MCFs, so if nonzero quadratic coefficients are present, they are just ignored.

**7.3.6 Member Data Documentation****7.3.6.1 int MCFCClass::MaxIter [protected]**

max number of iterations in which MCF Solver can find an optimal solution (0 = no limits)

**7.3.6.2 double MCFClass::MaxTime [protected]**

max time (in seconds) in which MCF Solver can find an optimal solution (0 = no limits)

**7.3.6.3 int MCFClass::status [protected]**

return status, see the comments to [MCFGetStatus\(\)](#) above. Note that the variable is defined int to allow derived classes to return their own specialized status codes

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

## 7.4 MCFClass::MCFException Class Reference

```
#include <MCFClass.h>
```

### Public Member Functions

- **MCFException** (const char \*const msg=0)
- const char \* **what** (void) const throw ()

#### 7.4.1 Detailed Description

Small class for exceptions. Derives from std::exception implementing the virtual method what() -- and since what is virtual, remember to always catch it by reference (catch exception &e) if you want the thing to work. [MCFException](#) class are thought to be of the "fatal" type, i.e., problems for which no solutions exists apart from aborting the program. Other kinds of exceptions can be properly handled by defining derived classes with more information.


The documentation for this class was generated from the following file:

- [MCFClass.h](#)

## 7.5 MCFSimplex Class Reference

```
#include <MCFSimplex.h>
```

Inheritance diagram for MCFSimplex:



### Classes

- struct **arcDType**
- struct **arcPType**
- struct **dualCandidType**
- struct **nodeDType**
- struct **nodePType**
- struct **primalCandidType**

### Public Types

- enum **SimplexParam** {  
     **kAlgPrimal** = **kLastParam**, **kAlgPricing**, **kNumCandList**, **kHotListSize**,  
     **kRecomputeFOLimits**, **kEpsOpt** }
- enum **enumPrcngRl** { **kDantzig** = 0, **kFirstEligibleArc**, **kCandidateListPivot** }

### Public Member Functions

- **MCFSimplex** (**cIndex** nmX=0, **cIndex** mmX=0)
- void **LoadNet** (**cIndex** nmX=0, **cIndex** mmX=0, **cIndex** pn=0, **cIndex** pm=0, **cFRow** pU=NULL, **cCRow** pC=NULL, **cFRow** pDfct=NULL, **cIndex\_Set** pSn=NULL, **cIndex\_Set** pEn=NULL)
- void **SetAlg** (bool UsPrml, char WhchPrc)
- void **SetPar** (int par, int val)
- void **SetPar** (int par, double val)
- void **SolveMCF** (void)
- void **MCFGetX** (register **FRow** F, register **Index\_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- void **MCFGetRC** (register **CRow** CR, register **cIndex\_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **CNumber** **MCFGetRC** (**cIndex** i)
- void **MCFGetPi** (register **CRow** P, register **cIndex\_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **FONumber** **MCFGetFO** (void)
- virtual void **MCFArcs** (register **Index\_Set** Startv, register **Index\_Set** Endv, register **cIndex\_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **Index** **MCFSNde** (**cIndex** i)
- **Index** **MCFENde** (**cIndex** i)

- void **MCFCosts** (register [CRow](#) Costv, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- [CNumber](#) **MCFCost** ([cIndex](#) i)
- void **MCFQCoef** (register [CRow](#) Qv, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- [CNumber](#) **MCFQCoef** ([cIndex](#) i)
- void **MCFUCaps** (register [FRow](#) UCapv, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- [FNumber](#) **MCFUCap** ([cIndex](#) i)
- void **MCFDfcts** (register [FRow](#) Dfctv, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- [FNumber](#) **MCFDfct** ([cIndex](#) i)
- void **ChgCosts** (register [cCRow](#) NCost, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- void **ChgCost** (register [Index](#) arc, [cCNumber](#) NCost)
- void **ChgQCoef** (register [cCRow](#) NQCoef=NULL, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- void **ChgQCoef** (register [Index](#) arc, [cCNumber](#) NQCoef)
- void **ChgDfcts** (register [cFRow](#) NDfct, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- void **ChgDfct** (register [Index](#) nod, [cFNumber](#) NDfct)
- void **ChgUCaps** (register [cFRow](#) NCap, register [cIndex\\_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- void **ChgUCap** (register [Index](#) arc, [cFNumber](#) NCap)
- void **CloseArc** ([cIndex](#) name)
- void **DelNode** ([cIndex](#) name)
- bool **IsClosedArc** ([cIndex](#) name)
- void **OpenArc** ([cIndex](#) name)
- [Index](#) **AddNode** ([cFNumber](#) aDfct)
- void **ChangeArc** ([cIndex](#) name, [cIndex](#) nSS=Inf< [Index](#) >(), [cIndex](#) nEN=Inf< [Index](#) >())
- void **DelArc** ([cIndex](#) name)
- [Index](#) **AddArc** ([cIndex](#) Start, [cIndex](#) End, [cFNumber](#) aU, [cCNumber](#) aC)
- bool **IsDeletedArc** ([cIndex](#) name)

### 7.5.1 Detailed Description

The [MCFSimplex](#) class derives from the abstract base class [MCFClass](#), thus sharing its (standard) interface, and implements both the Primal and Dual network simplex algorithms for solving (Linear and Quadratic) Min Cost Flow problems

### 7.5.2 Member Enumeration Documentation

#### 7.5.2.1 enum [MCFSimplex::enumPrcngRI](#)

Public enum describing the pricing rules in [MCFSimplex::SetAlg\(\)](#).

**Enumerator:**

- kDantzig*** Dantzig's rule (most violated constraint).
- kFirstEligibleArc*** First eligible arc in round-robin.
- kCandidateListPivot*** Candidate List Pivot Rule.

### 7.5.2.2 enum MCFSimplex::SimplexParam

Public enum describing the parameters of [MCFSimplex](#).

#### Enumerator:

*kAlgPrimal* `kAlgPrimal == MCFCClass::kLastParam == 6` parameter to set algorithm (Primal/Dual):

*kAlgPricing* parameter to set algorithm of pricing

*kNumCandList* list for Candidate List Pivot method parameter to set the number of candidate

*kHotListSize* for Candidate List Pivot method parameter to set the size of Hot List

*kRecomputeFOLimits* parameter to set the number of iterations in which quadratic Primal Simplex computes "manually" the f.o. value

*kEpsOpt* parameter to set the precision of the objective function value for the quadratic case of the Primal Simplex

## 7.5.3 Constructor & Destructor Documentation

### 7.5.3.1 MCFSimplex::MCFSimplex (cIndex *nm*x = 0, cIndex *mm*x = 0)

Constructor of the class, as in [MCFCClass::MCFCClass\(\)](#).

## 7.5.4 Member Function Documentation

### 7.5.4.1 MCFSimplex::Index MCFSimplex::AddArc (cIndex *Start*, cIndex *End*, cFNumber *aU*, cCNumber *aC*) [**virtual**]

Add the new arc ( *Start* , *End* ) with cost *aC* and capacity *aU*, returning its name. `Inf<Index>()` is returned if there is no room for a new arc. Remember that arc names go from 0 to *mmax* - 1.

Implements [MCFCClass](#).

### 7.5.4.2 MCFSimplex::Index MCFSimplex::AddNode (cFNumber *aDfct*) [**virtual**]

Add a new node with deficit *aDfct*, returning its name. `Inf<Index>()` is returned if there is no room for a new node. Remember that the node names are either { 0 .. *nmax* - 1 } or { 1 .. *nmax* }, depending on the value of `USENAME0`.

Implements [MCFCClass](#).

### 7.5.4.3 void MCFSimplex::ChangeArc (cIndex *name*, cIndex *nSN* = `Inf< Index >()`, cIndex *nEN* = `Inf< Index >()`) [**virtual**]

Change the starting and/or ending node of arc 'name' to *nSN* and *nEN*. Each parameter being `Inf<Index>()` means to leave the previous starting or ending node untouched. When this method is called 'name' can be either the name of a "normal" arc or that of a "closed" arc [see [CloseArc\(\)](#) above]: in the latter case, at the end of [ChangeArc\(\)](#) the arc is \*still closed\*, and it remains so until [OpenArc\( name \)](#) [see above] is called.

Implements [MCFCClass](#).

**7.5.4.4 void MCFSimplex::CloseArc (cIndex *name*) [virtual]**

"Close" the arc '*name*'. Although all the associated information (name, cost, capacity, end and start node) is kept, the arc is removed from the problem until `OpenArc(i)` [see below] is called.

"closed" arcs always have 0 flow, but are otherwise counted as any other arc; for instance, `MCFm()` does *not* decrease as an effect of a call to `CloseArc()`. How this closure is implemented is solver-specific.

Implements [MCFClass](#).

**7.5.4.5 void MCFSimplex::DelArc (cIndex *name*) [virtual]**

Delete the arc '*name*'. Unlike "closed" arcs, all the information associated with a deleted arc is lost and '*name*' is made available as a name for new arcs to be created with `AddArc()` [see below].

If furthermore '*name*' is the last arc, the number of arcs as reported by `MCFm()` is reduced by at least one, until the *m*-th arc is not a deleted one. Otherwise, the flow on the arc is always ensured to be 0.

Implements [MCFClass](#).

**7.5.4.6 void MCFSimplex::DelNode (cIndex *name*) [virtual]**

Delete the node '*name*'.

For any value of '*name*', all incident arcs to that node are closed [see `CloseArc()` above] (*not* Deleted, see `DelArc()` below) and the deficit is set to zero.

If furthermore '*name*' is the last node, the number of nodes as reported by `MCFn()` is reduced by at least one, until the *n*-th node is not a deleted one.

Implements [MCFClass](#).

**7.5.4.7 bool MCFSimplex::IsClosedArc (cIndex *name*) [virtual]**

`IsClosedArc()` returns true if and only if the arc '*name*' is closed.

Implements [MCFClass](#).

**7.5.4.8 bool MCFSimplex::IsDeletedArc (cIndex *name*) [virtual]**

Return true if and only if the arc '*name*' is deleted. It should only be called with `name < MCFm()`, as every other arc is deleted by definition.

Implements [MCFClass](#).

**7.5.4.9 void MCFSimplex::LoadNet (cIndex *nm*x = 0, cIndex *mm*x = 0, cIndex *pn* = 0, cIndex *pm* = 0, cFRow *pU* = NULL, cCRow *pC* = NULL, cFRow *pDfct* = NULL, cIndex\_Set *pSn* = NULL, cIndex\_Set *pEn* = NULL) [virtual]**

Inputs a new network, as in `MCFClass::LoadNet()`.

Implements [MCFClass](#).

**7.5.4.10 MCFClass::CNumber MCFSimplex::MCFCost (cIndex *i*) [inline, virtual]**

Return the cost of the *i*-th arc.

Implements [MCFClass](#).

**7.5.4.11 MCFClass::FNumber MCFSimplex::MCFDfct (cIndex *i*) [inline, virtual]**

Return the deficit of the *i*-th node.

**Note**

Here node "names" go from 0 to *n* - 1, regardless to the value of USERNAME0; hence, if USERNAME0 == 0 then the first node is "named 1", but its deficit is returned by MCFDfct( 0 ).

Implements [MCFClass](#).

**7.5.4.12 MCFClass::Index MCFSimplex::MCFENde (cIndex *i*) [inline, virtual]**

Return the ending (head) node of the arc '*i*'.

**Note**

If USERNAME0 == 0 then the returned node names will be in the range 1 .. *n*, while if USERNAME0 == 1 the returned node names will be in the range 0 .. *n* - 1.

Implements [MCFClass](#).

**7.5.4.13 MCFSimplex::FNumber MCFSimplex::MCFGetFO (void) [virtual]**

Return the objective function value of the primal solution currently returned by [MCFGetX\(\)](#).

If [MCFGetStatus\(\)](#) == kOK, this is guaranteed to be the optimal objective function value of the problem (to within the optimality tolerances), but only prior to any call to [HaveNewX\(\)](#) that returns true. [MCFGetFO\(\)](#) typically returns Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnfeasible and

- Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnbounded. If [MCFGetStatus\(\)](#) == kStopped and [MCFGetFO\(\)](#) returns a finite value, it must be an upper bound on the optimal objective function value (typically, the objective function value of one primal feasible solution).

Implements [MCFClass](#).

**7.5.4.14 MCFSimplex::CNumber MCFSimplex::MCFGetRC (cIndex *i*) [virtual]**

Return the reduced cost of the *i*-th arc. This information should be cheaply available in most implementations.

**Note**

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Implements [MCFClass](#).



**7.5.4.15 MCFClass::Index MCFSimplex::MCFNde (cIndex *i*) [inline, virtual]**

Return the starting (tail) node of the arc '*i*'.

**Note**

If USENAME0 == 0 then the returned node names will be in the range 1 .. *n*, while if USENAME0 == 1 the returned node names will be in the range 0 .. *n* - 1.

Implements [MCFClass](#).

**7.5.4.16 MCFClass::FNumber MCFSimplex::MCFUCap (cIndex *i*) [inline, virtual]**

Return the capacity of the *i*-th arc.

Implements [MCFClass](#).

**7.5.4.17 void MCFSimplex::OpenArc (cIndex *name*) [virtual]**

Restore the previously closed arc '*name*'. It is an error to open an arc that has not been previously closed.

Implements [MCFClass](#).

**7.5.4.18 void MCFSimplex::SetAlg (bool *UsPrml*, char *WhchPrc*)**

Select which algorithm (Primal vs Dual Network Simplex), and which pricing rule within the algorithm, is used.

If *UsPrml* == TRUE then the Primal Network Simplex algorithm is used, otherwise the Dual Network Simplex is used.

The allowed values for *WhchPrc* are:

- *kDantzig* Dantzig's pricing rule, i.e., most violated dual constraint; this can only be used with the Primal Network Simplex
- *kFirstEligibleArcA* "dumb" rule, first eligible arc in round-robin;
- *kCandidateListPivot* Candidate List Pivot Rule

If this method is *\*not\** called, the Primal Network Simplex with the Candidate List Pivot Rule (the best setting on most instances) is used.

**7.5.4.19 void MCFSimplex::SetPar (int *par*, double *val*) [virtual]**

Set general float parameters.

**Parameters**

*par* is parameter to set; enum SimplexParam (in addition to MCFParam) can be used, but this method accepts an int value (every enum is an int value), so it can be extended by derived classes for the setting of their parameters

*value* is the value to assign to the parameter.

Reimplemented from [MCFClass](#).

**7.5.4.20 void MCFSimplex::SetPar (int *par*, int *val*) [virtual]**

Set general integer parameters.

**Parameters**

*par* is parameter to set; enum SimplexParam (in addition to MCFParam) can be used, but this method accepts an int value (every enum is an int value), so it can be extended by derived classes for the setting of their parameters

*value* is the value to assign to the parameter.

Reimplemented from [MCFClass](#).

**7.5.4.21 void MCFSimplex::SolveMCF (void) [virtual]**

Solver of the Min Cost Flow Problem. Attempts to solve the MCF instance currently loaded in the object.

Implements [MCFClass](#).

The documentation for this class was generated from the following files:

- [MCFSimplex.h](#)
- MCFSimplex.C

## 7.6 MCFClass::MCFState Class Reference

```
#include <MCFClass.h>
```

### 7.6.1 Detailed Description

Base class for representing the internal state of the MCF algorithm.

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

## 7.7 OPTrand Class Reference

```
#include <OPTUtils.h>
```

### Public Member Functions

- [OPTrand](#) (void)  
*constructor of the class*
- double [rand](#) (void)
- void [srand](#) (long seed)  
*Seeds the random generator for this instance of [OPTrand](#).*

### 7.7.1 Detailed Description

Provide a common interface to the different random generators that are available in different platforms.

### 7.7.2 Member Function Documentation

#### 7.7.2.1 double OPTrand::rand (void) [inline]

Returns a random number uniformly distributed in [0, 1).

#### Note

each object of class [OPTrand](#) has its own sequence, so that multiple [OPTrand](#) objects being used within the same program do not interfere with each other (as opposed to what C random routines would do).

The documentation for this class was generated from the following file:

- [OPTUtils.h](#)

# Chapter 8

## File Documentation

### 8.1 Main.C File Reference

```
#include <fstream>
#include <sstream>
#include "MCFSimplex.h"
#include <stdlib.h>
#include <time.h>
#include <iostream>
#include <iomanip>
```

#### Defines

- `#define MCF_SOLVER MCFSimplex`
- `#define PRINT_RESULTS 0`

#### Functions

- `template<class T >`  
  `T ABS (const T x)`
- `void SkipComments (ifstream &iParam, string &buf)`
- `void SetParam (MCFClass *mcf)`
- `int main (int argc, char **argv)`

#### 8.1.1 Detailed Description

Sample Main file to illustrate the use of any solver deriving from [MCFClass](#). By changing just \*two lines of code\* and little more (see comment PECULIARITY, if exists) the file works with any derived solver.

An instance of a Min Cost Flow problem in DIMACS standard format is read from file and solved. In addition, the same problem can be written on a file in MPS format.

**Version**

4.00

**Date**

30 - 12 - 2009

**Author**

Alessandro Bertolini  
Operations Research Group  
Dipartimento di Informatica  
Universita' di Pisa  
Antonio Frangioni  
Operations Research Group  
Dipartimento di Informatica  
Universita' di Pisa

## 8.2 MCFClass.h File Reference

```
#include "OPTUtils.h"  
#include <iomanip>  
#include <sstream>
```

### Classes

- class [MCFClass](#)
- class [MCFClass::Inf< T >](#)
- class [MCFClass::Eps< T >](#)
- class [MCFClass::MCException](#)
- class [MCFClass::MCState](#)

### Defines

- #define [USENAME0](#) 0

### 8.2.1 Detailed Description

Header file for the abstract base class [MCFClass](#), which defines a standard interface for (linear or convex quadratic separable) Min Cost Flow Problem solvers, to be implemented as derived classes.

#### Version

3.00

#### Date

30 - 12 - 2009

#### Author

Alessandro Bertolini  
Operations Research Group  
Dipartimento di Informatica  
Universita' di Pisa  
Antonio Frangioni  
Operations Research Group  
Dipartimento di Informatica  
Universita' di Pisa  
Claudio Gentile  
Istituto di Analisi di Sistemi e Informatica  
Consiglio Nazionale delle Ricerche

Copyright &copy; 1996 - 2009 by Antonio Frangioni, Claudio Gentile

## 8.3 MCFSimplex.h File Reference

```
#include "MCFCClass.h"
```

### Classes

- class [MCFSimplex](#)
- struct **MCFSimplex::nodePType**
- struct **MCFSimplex::nodeDType**
- struct **MCFSimplex::arcPType**
- struct **MCFSimplex::arcDType**
- struct **MCFSimplex::primalCandidType**
- struct **MCFSimplex::dualCandidType**

### Defines

- #define [QUADRATICCOST](#) 0

### 8.3.1 Detailed Description

Linear Min Cost Flow problems solver, code by Alessandro Bertolini. Conforms to the standard MCF interface defined in [MCFCClass.h](#).

1.0

#### Date

01 - 03 - 2008

#### Author

(original C++ code)  
Alessandro Bertolini



## 8.4 OPTUtils.h File Reference

```
#include <sys/timeb.h>
#include <limits.h>
#include <iostream>
#include <sstream>
```

### Classes

- class [OPTrand](#)

### Defines

- #define [OPT\\_USE\\_NAMESPACES](#) 0
- #define [OPT\\_TIMERS](#) 5
- #define [OPT\\_RANDOM](#) 1
- #define **NULL** 0
- #define **CLK\_TCK** 100

### Functions

- template<class T >  
void [DfltDfInpt](#) (istream \*iStrm, T &Param, const T Dflt, const char cmntc= '#')

### 8.4.1 Detailed Description

Small classes are provided for:

- reading the time of a code;
- generating random numbers.

The classes can be adapted to different environments setting a compile-time switch in this file.

Additionally, a function is provided for safely reading parameters out of a stream.

#### Version

1.00

#### Date

04 - 10 - 2010

#### Author

Antonio Frangioni  
Operations Research Group  
Dipartimento di Informatica  
Universita' di Pisa

Copyright &copy; 1994 - 2010 by Antonio Frangioni

# Index

- ~MCFClass
  - MCFClass, [27](#)
- AddArc
  - MCFClass, [28](#)
  - MCFSimplex, [50](#)
- AddNode
  - MCFClass, [28](#)
  - MCFSimplex, [50](#)
- ChangeArc
  - MCFClass, [28](#)
  - MCFSimplex, [50](#)
- CheckDSol
  - MCFClass, [28](#)
- CheckPSol
  - MCFClass, [28](#)
- ChgCost
  - MCFClass, [28](#)
- ChgCosts
  - MCFClass, [29](#)
- ChgDfct
  - MCFClass, [29](#)
- ChgDfcts
  - MCFClass, [29](#)
- ChgQCoef
  - MCFClass, [30](#)
- ChgUCap
  - MCFClass, [30](#)
- ChgUCaps
  - MCFClass, [30](#)
- Classes in MCFClass.h, [10](#)
- Classes in MCFSimplex.h, [12](#)
- Classes in OPTutils.h, [15](#)
- CloseArc
  - MCFClass, [31](#)
  - MCFSimplex, [50](#)
- Compile-time switches in MCFClass.h, [9](#)
- Compile-time switches in MCFSimplex.h, [11](#)
- Compile-time switches in OPTutils.h, [13](#)
- DelArc
  - MCFClass, [31](#)
  - MCFSimplex, [51](#)
- DelNode
  - MCFClass, [31](#)
  - MCFSimplex, [51](#)
- DfltDfInpt
  - OPTtypes\_FUNCTIONS, [16](#)
- enumPrcngRl
  - MCFSimplex, [49](#)
- ETZ
  - MCFClass, [31](#)
- FONumber
  - MCFClass, [26](#)
- Functions in OPTutils.h, [16](#)
- GetPar
  - MCFClass, [31](#), [32](#)
- GEZ
  - MCFClass, [32](#)
- GT
  - MCFClass, [32](#)
- GTZ
  - MCFClass, [32](#)
- HaveNewPi
  - MCFClass, [32](#)
- HaveNewX
  - MCFClass, [32](#)
- IsClosedArc
  - MCFClass, [33](#)
  - MCFSimplex, [51](#)
- IsDeletedArc
  - MCFClass, [33](#)
  - MCFSimplex, [51](#)
- kAlgPricing
  - MCFSimplex, [50](#)
- kAlgPrimal
  - MCFSimplex, [50](#)
- kCandidateListPivot
  - MCFSimplex, [49](#)
- kDantzig
  - MCFSimplex, [49](#)
- kDimacs
  - MCFClass, [26](#)
- kEpsCst

- MCFCClass, 27
- kEpsDfct
  - MCFCClass, 27
- kEpsFlw
  - MCFCClass, 27
- kEpsOpt
  - MCFSimplex, 50
- kError
  - MCFCClass, 27
- kFirstEligibleArc
  - MCFSimplex, 49
- kFWMPS
  - MCFCClass, 26
- kHotListSize
  - MCFSimplex, 50
- kLastParam
  - MCFCClass, 27
- kMaxIter
  - MCFCClass, 26
- kMaxTime
  - MCFCClass, 26
- kMPS
  - MCFCClass, 26
- kNo
  - MCFCClass, 26
- kNumCandList
  - MCFSimplex, 50
- kOK
  - MCFCClass, 27
- kQDimacs
  - MCFCClass, 26
- kRecomputeFOLimits
  - MCFSimplex, 50
- kReopt
  - MCFCClass, 27
- kStopped
  - MCFCClass, 27
- kUnbounded
  - MCFCClass, 27
- kUnfeasible
  - MCFCClass, 27
- kUnSolved
  - MCFCClass, 27
- kYes
  - MCFCClass, 26
- LEZ
  - MCFCClass, 33
- LoadDMX
  - MCFCClass, 33
- LoadNet
  - MCFCClass, 34
  - MCFSimplex, 51
- LT
  - MCFCClass, 35
- LTZ
  - MCFCClass, 35
- Main.C, 57
- MaxIter
  - MCFCClass, 45
- MaxTime
  - MCFCClass, 45
- MCFAnswer
  - MCFCClass, 26
- MCFArcs
  - MCFCClass, 35
- MCFCClass, 19
  - ~MCFCClass, 27
  - AddArc, 28
  - AddNode, 28
  - ChangeArc, 28
  - CheckDSol, 28
  - CheckPSol, 28
  - ChgCost, 28
  - ChgCosts, 29
  - ChgDfct, 29
  - ChgDfcts, 29
  - ChgQCoef, 30
  - ChgUCap, 30
  - ChgUCaps, 30
  - CloseArc, 31
  - DelArc, 31
  - DelNode, 31
  - ETZ, 31
  - FONumber, 26
  - GetPar, 31, 32
  - GEZ, 32
  - GT, 32
  - GTZ, 32
  - HaveNewPi, 32
  - HaveNewX, 32
  - IsClosedArc, 33
  - IsDeletedArc, 33
  - kDimacs, 26
  - kEpsCst, 27
  - kEpsDfct, 27
  - kEpsFlw, 27
  - kError, 27
  - kFWMPS, 26
  - kLastParam, 27
  - kMaxIter, 26
  - kMaxTime, 26
  - kMPS, 26
  - kNo, 26
  - kOK, 27
  - kQDimacs, 26
  - kReopt, 27

- kStopped, 27
- kUnbounded, 27
- kUnfeasible, 27
- kUnsolved, 27
- kYes, 26
- LEZ, 33
- LoadDMX, 33
- LoadNet, 34
- LT, 35
- LTZ, 35
- MaxIter, 45
- MaxTime, 45
- MCFAnswer, 26
- MCFArcs, 35
- MCFClass, 27
- MFCFCost, 35
- MFCFCosts, 35
- MCFDfct, 36
- MCFDfcts, 36
- MCFENde, 36
- MCFENdes, 37
- MCFFIfrmt, 26
- MCFGetDFO, 37
- MCFGetFO, 37
- MCFGetPi, 37
- MCFGetRC, 38
- MCFGetState, 38
- MCFGetStatus, 39
- MCFGetUnbCycl, 39
- MCFGetUnfCut, 40
- MCFGetX, 40
- MCFm, 40
- MCFmmax, 40
- MCFn, 41
- MCFnmax, 41
- MCFParam, 26
- MCFPutState, 41
- MCFQCoef, 41, 42
- MCFSNde, 42
- MCFSNdes, 42
- MCFStatus, 27
- MCFUCap, 42
- MCFUCaps, 42, 43
- OpenArc, 43
- PreProcess, 43
- SetMCFTIME, 43
- SetPar, 43, 44
- SolveMCF, 44
- status, 46
- TimeMCF, 45
- WriteMCF, 45
- MCFClass.h, 59
- MCFClass::Eps, 17
- MCFClass::Inf, 18
- MCFClass::MCFException, 47
- MCFClass::MCFState, 55
- MCFCLASS\_MACROS
  - USENAME0, 9
- MFCFCost
  - MCFClass, 35
  - MCFSimplex, 51
- MFCFCosts
  - MCFClass, 35
- MCFDfct
  - MCFClass, 36
  - MCFSimplex, 52
- MCFDfcts
  - MCFClass, 36
- MCFENde
  - MCFClass, 36
  - MCFSimplex, 52
- MCFENdes
  - MCFClass, 37
- MCFFIfrmt
  - MCFClass, 26
- MCFGetDFO
  - MCFClass, 37
- MCFGetFO
  - MCFClass, 37
  - MCFSimplex, 52
- MCFGetPi
  - MCFClass, 37
- MCFGetRC
  - MCFClass, 38
  - MCFSimplex, 52
- MCFGetState
  - MCFClass, 38
- MCFGetStatus
  - MCFClass, 39
- MCFGetUnbCycl
  - MCFClass, 39
- MCFGetUnfCut
  - MCFClass, 40
- MCFGetX
  - MCFClass, 40
- MCFm
  - MCFClass, 40
- MCFmmax
  - MCFClass, 40
- MCFn
  - MCFClass, 41
- MCFnmax
  - MCFClass, 41
- MCFParam
  - MCFClass, 26
- MCFPutState
  - MCFClass, 41
- MCFQCoef

- MCFClass, [41, 42](#)
- MCFSimplex, [48](#)
  - AddArc, [50](#)
  - AddNode, [50](#)
  - ChangeArc, [50](#)
  - CloseArc, [50](#)
  - DelArc, [51](#)
  - DelNode, [51](#)
  - enumPrngRl, [49](#)
  - IsClosedArc, [51](#)
  - IsDeletedArc, [51](#)
  - kAlgPricing, [50](#)
  - kAlgPrimal, [50](#)
  - kCandidateListPivot, [49](#)
  - kDantzig, [49](#)
  - kEpsOpt, [50](#)
  - kFirstEligibleArc, [49](#)
  - kHotListSize, [50](#)
  - kNumCandList, [50](#)
  - kRecomputeFOLimits, [50](#)
  - LoadNet, [51](#)
  - MCFCost, [51](#)
  - MCFDfct, [52](#)
  - MCFENde, [52](#)
  - MCFGetFO, [52](#)
  - MCFGetRC, [52](#)
  - MCFSimplex, [50](#)
  - MCFSNde, [52](#)
  - MCFUCap, [53](#)
  - OpenArc, [53](#)
  - SetAlg, [53](#)
  - SetPar, [53](#)
  - SimplexParam, [49](#)
  - SolveMCF, [54](#)
- MCFSimplex.h, [60](#)
- MCFSimplex\_MACROS
  - QUADRATICCOST, [11](#)
- MCFSNde
  - MCFClass, [42](#)
  - MCFSimplex, [52](#)
- MCFSNdes
  - MCFClass, [42](#)
- MCFStatus
  - MCFClass, [27](#)
- MCFUCap
  - MCFClass, [42](#)
  - MCFSimplex, [53](#)
- MCFUCaps
  - MCFClass, [42, 43](#)
- OpenArc
  - MCFClass, [43](#)
  - MCFSimplex, [53](#)
- OPT\_RANDOM
  - OPTUTILS\_MACROS, [13](#)
- OPT\_TIMERS
  - OPTUTILS\_MACROS, [13](#)
- OPT\_USE\_NAMESPACES
  - OPTUTILS\_MACROS, [14](#)
- OPTrand, [56](#)
  - rand, [56](#)
- OPTtypes\_FUNCTIONS
  - DfltDfInpt, [16](#)
- OPTUtils.h, [61](#)
- OPTUTILS\_MACROS
  - OPT\_RANDOM, [13](#)
  - OPT\_TIMERS, [13](#)
  - OPT\_USE\_NAMESPACES, [14](#)
- PreProcess
  - MCFClass, [43](#)
- QUADRATICCOST
  - MCFSimplex\_MACROS, [11](#)
- rand
  - OPTrand, [56](#)
- SetAlg
  - MCFSimplex, [53](#)
- SetMCFTime
  - MCFClass, [43](#)
- SetPar
  - MCFClass, [43, 44](#)
  - MCFSimplex, [53](#)
- SimplexParam
  - MCFSimplex, [49](#)
- SolveMCF
  - MCFClass, [44](#)
  - MCFSimplex, [54](#)
- status
  - MCFClass, [46](#)
- TimeMCF
  - MCFClass, [45](#)
- USERNAME0
  - MCFCCLASS\_MACROS, [9](#)
- WriteMCF
  - MCFClass, [45](#)